

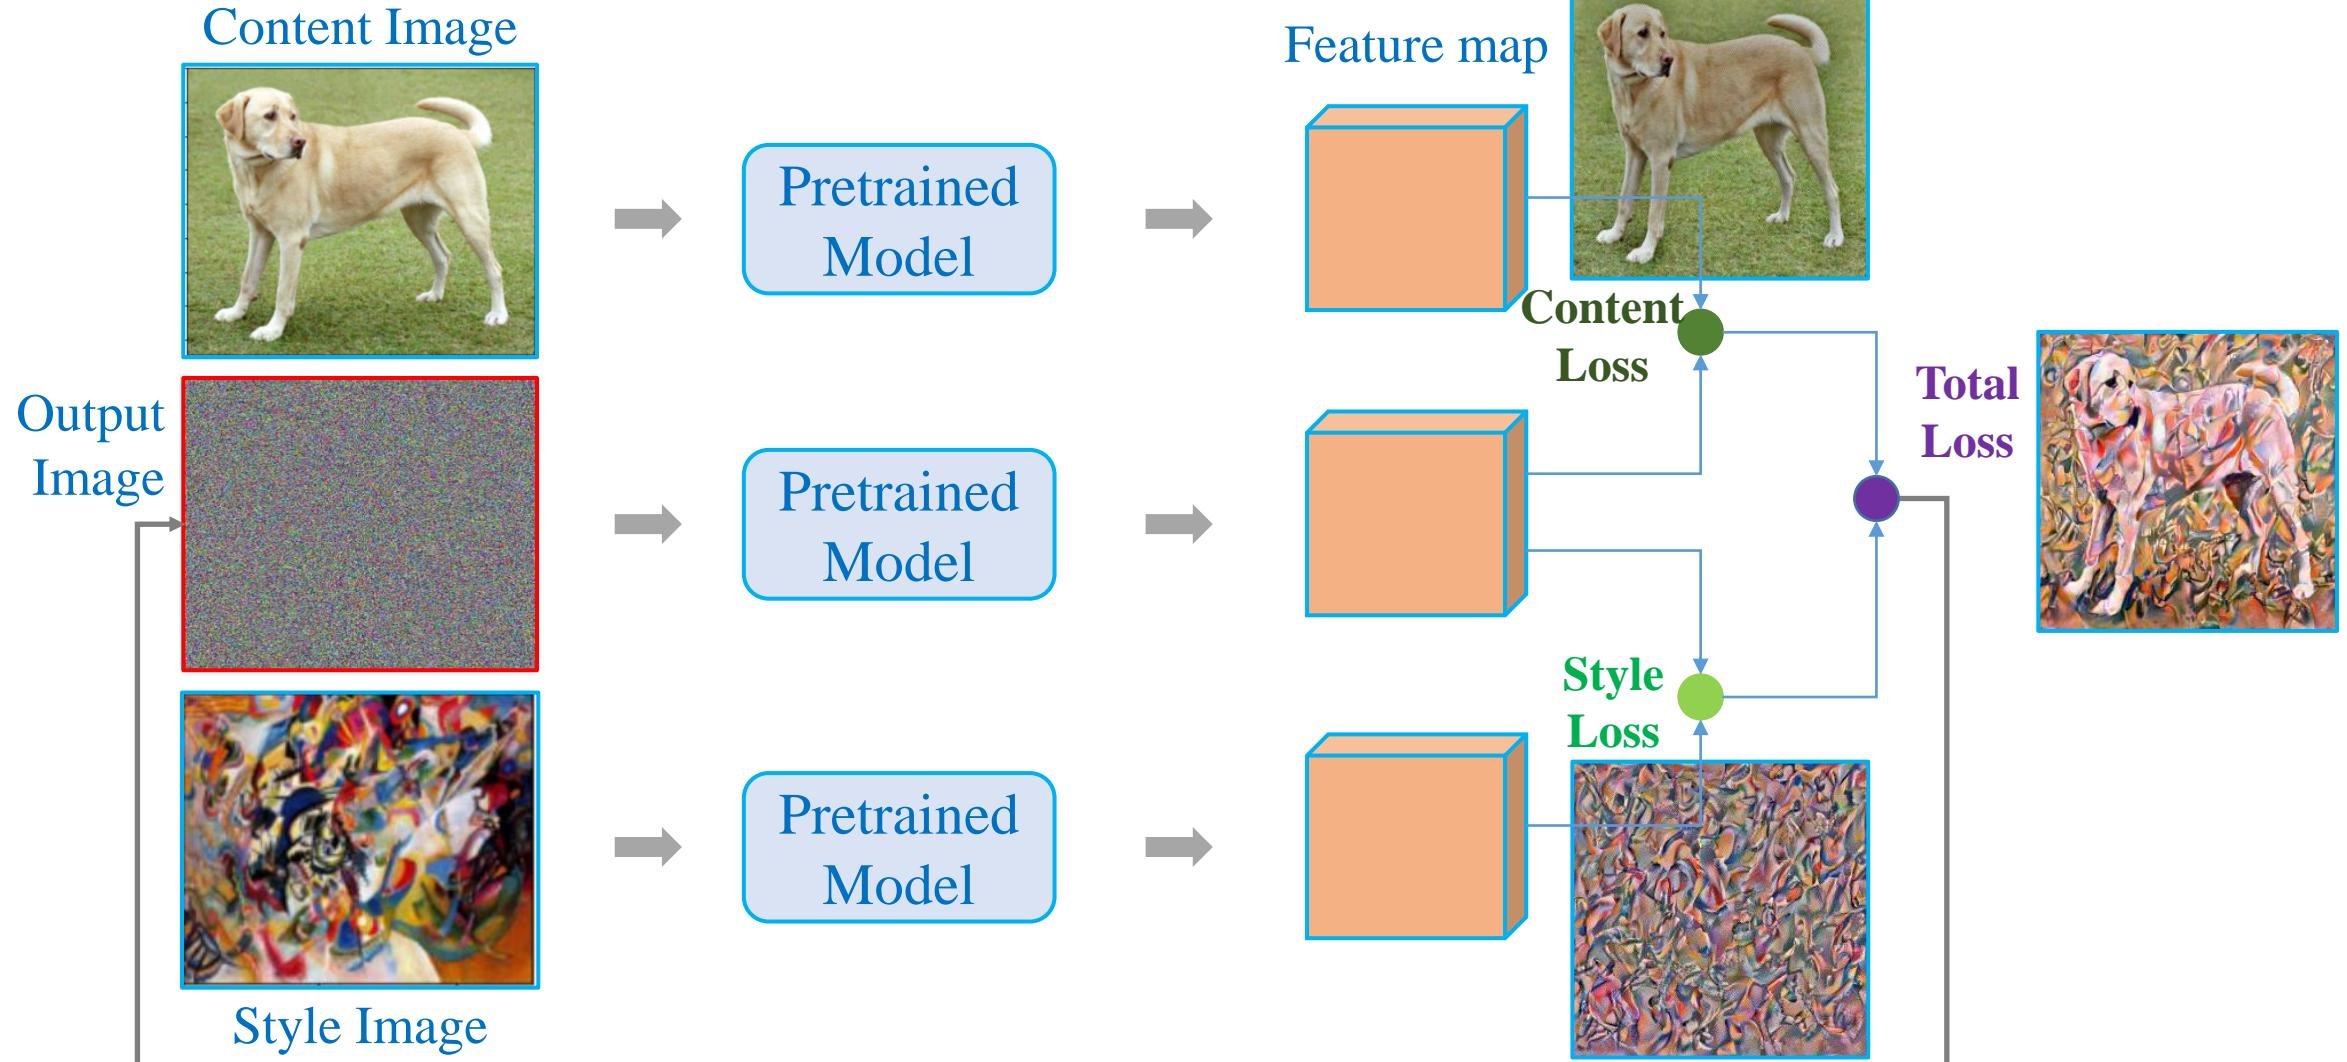
Style Transfer

Beyond the content and style losses

Quang-Vinh Dinh
Ph.D. in Computer Science

Objectives

- ✓ Study about content loss + style loss
- ✓ Implementation of a simple method



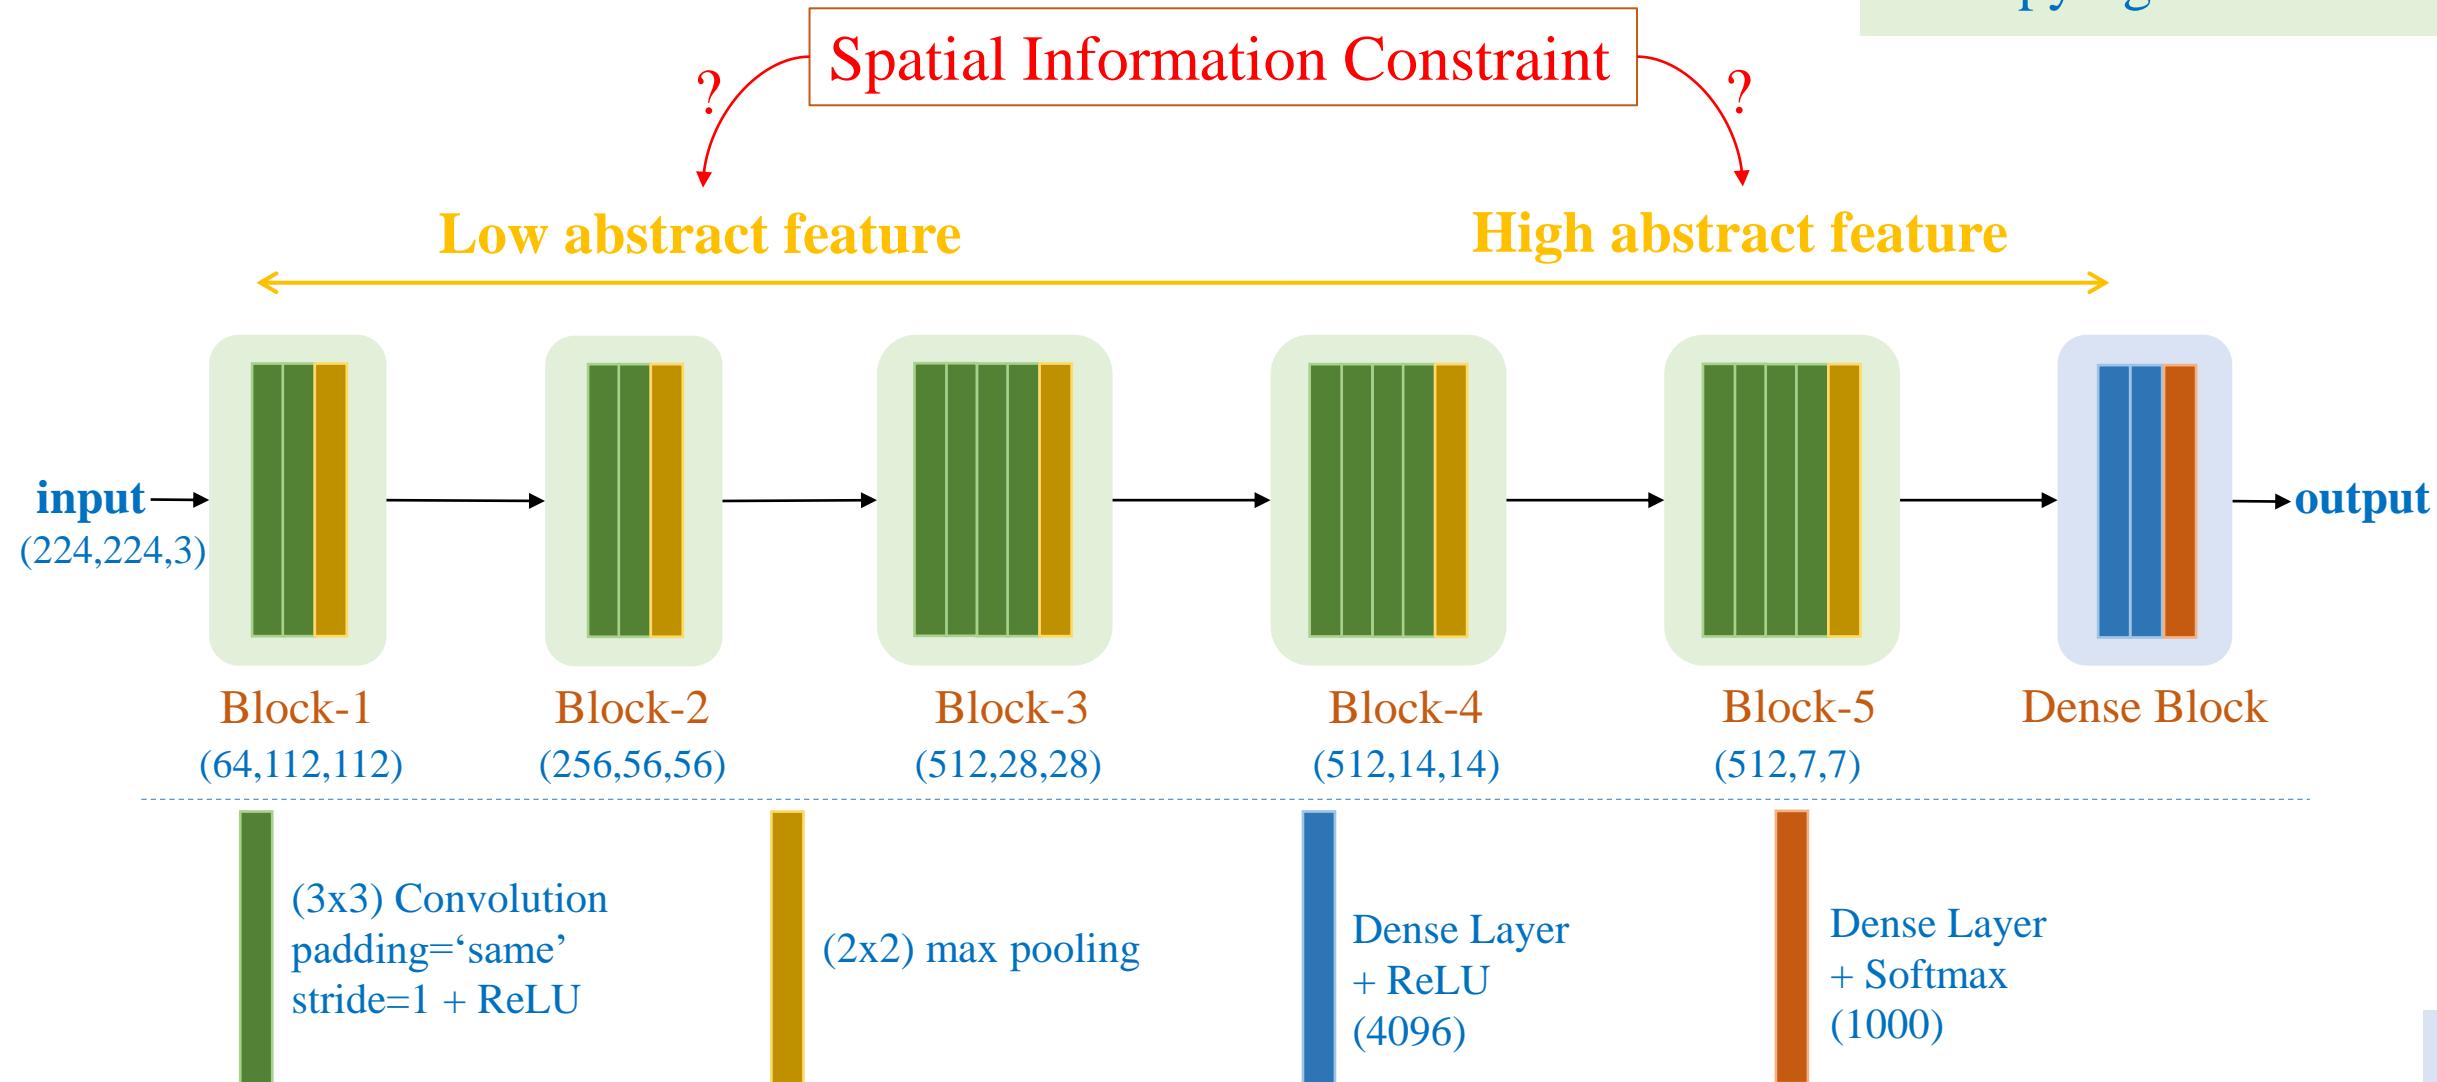
Outline

- Quick Review
- Content Loss+Style Loss
- Implementation
- Einstein Summation
- Extensions

Review

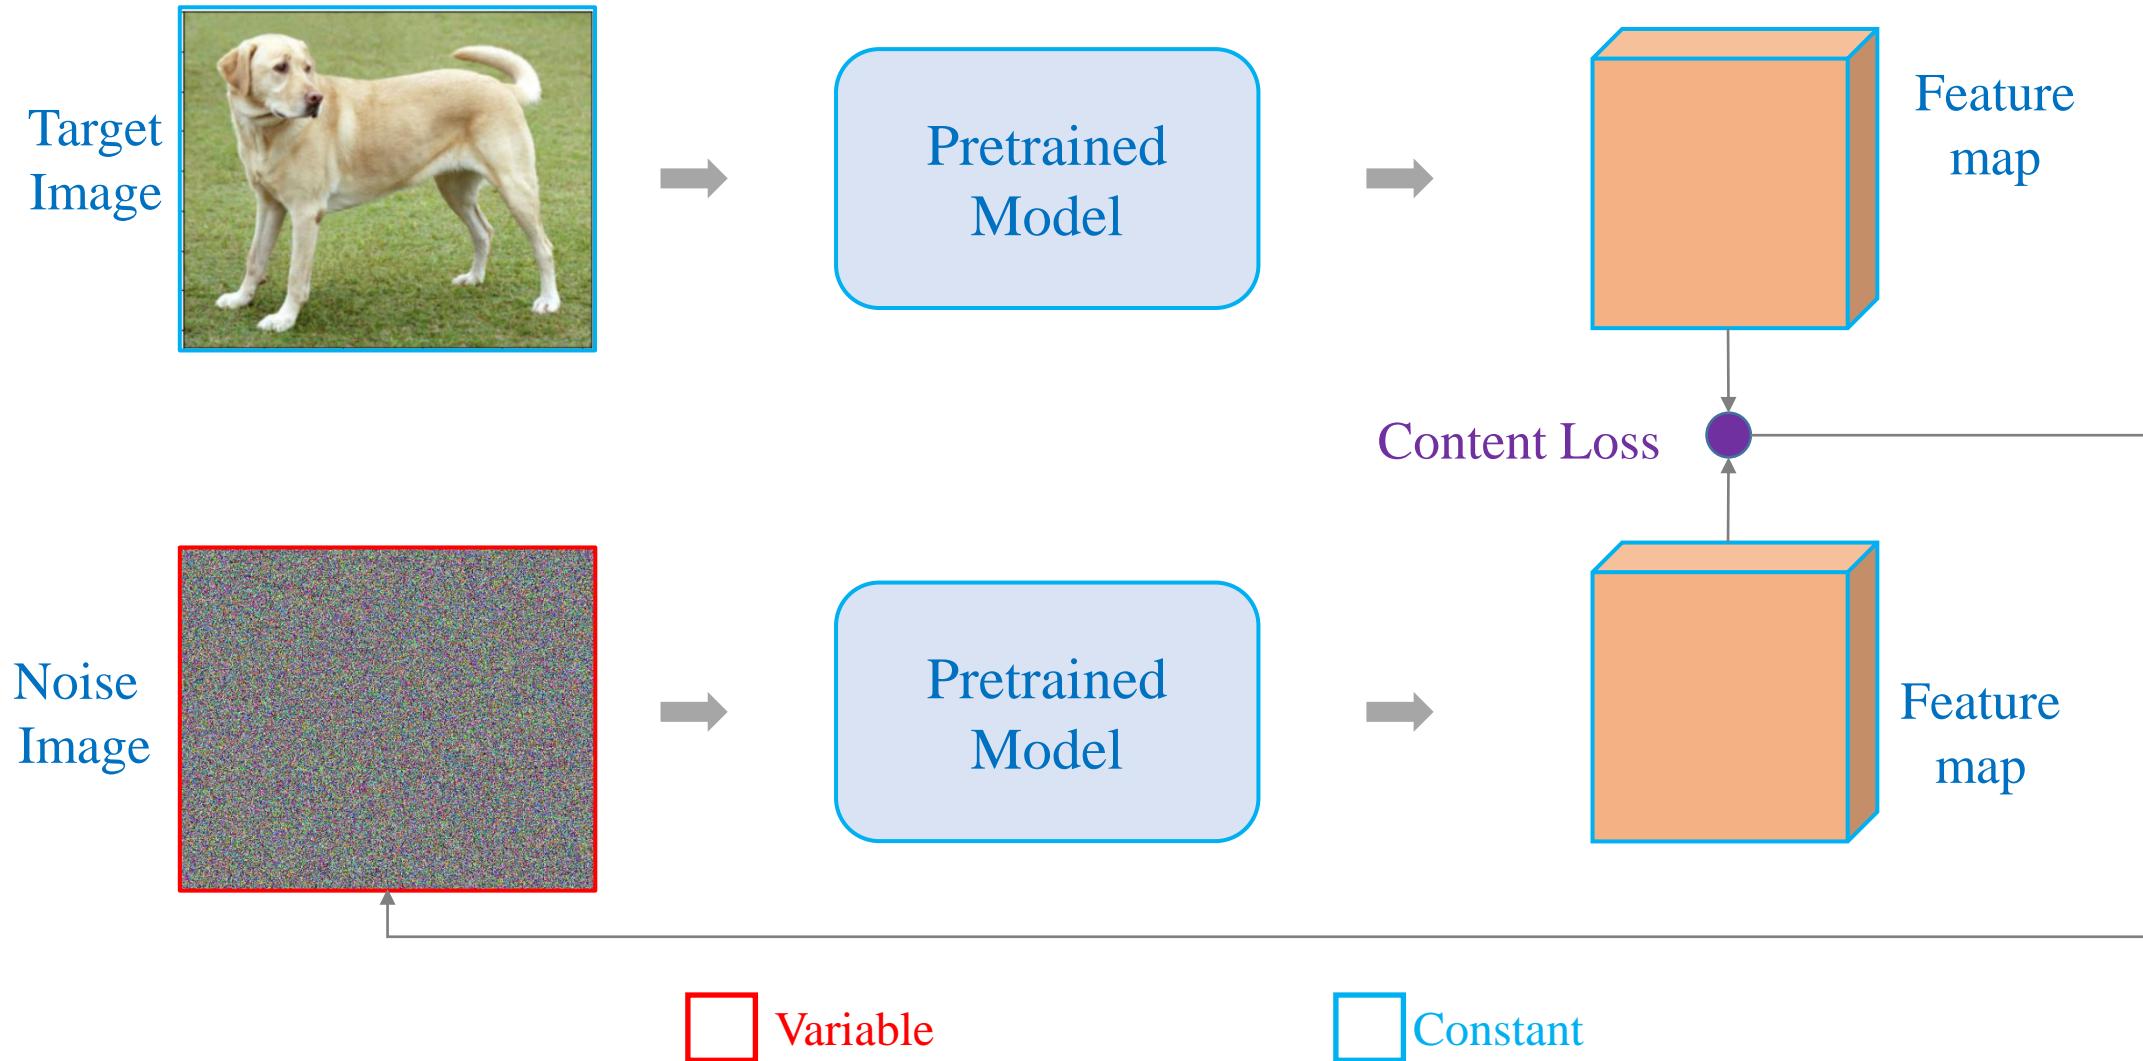
Local and Global
Detection Problem
Copying Problem

❖ Pretrained model (VGG19)



Review

❖ Content Loss



Review

❖ Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

```
import torch

x = torch.Tensor([[4, 0, 3, 1],
                  [3, 0, 1, 2],
                  [4, 5, 0, 3]])
G = torch.mm(x, x.t())
print(G)

tensor([[26., 17., 19.],
        [17., 14., 18.],
        [19., 18., 50.]])
```

The diagram illustrates the computation of the Gram matrix. It shows two 4x4 matrices, V and V^T , being multiplied to yield the Gram matrix. Matrix V has columns with values [4, 3, 1] and [0, 0, 5]. Matrix V^T has rows with values [4, 3, 4], [0, 0, 5], [3, 1, 0], and [1, 2, 3]. The resulting Gram matrix is a 4x4 matrix with diagonal elements 26, 17, 19, 17, 14, 18, 19, 18, 50.

The diagram illustrates the computation of the Gram matrix. It shows two 4x4 matrices, V and V^T , being multiplied to yield the Gram matrix. Matrix V has columns with values [4, 3, 1] and [0, 0, 5]. Matrix V^T has rows with values [4, 3, 4], [0, 0, 5], [3, 1, 0], and [1, 2, 3]. The resulting Gram matrix is a 4x4 matrix with diagonal elements 26, 17, 19, 17, 14, 18, 19, 18, 50.

Review

❖ Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

```
import torch

x = torch.Tensor([[4, 0, 3, 1],
                  [3, 0, 1, 2],
                  [4, 5, 0, 3]])
G = torch.einsum('ij,jk->ik', x, x.t())
print(G)

tensor([[26., 17., 19.],
        [17., 14., 18.],
        [19., 18., 50.]])
```

The diagram shows two 3x4 matrices, V and V^T , being multiplied to produce the Gram matrix. V has columns colored green, blue, green, and yellow. V^T has rows colored green, blue, and yellow. The result is a 3x3 matrix where each element is the dot product of a row from V and a column from V^T .

4	0	3	1
3	0	1	2
4	5	0	3

\bullet

4	3	4
0	0	5
3	1	0

$=$

26	17	19
17	14	18
19	18	50

Gram matrix

The diagram shows two 3x4 matrices, V and V^T , being multiplied to produce the Gram matrix. V has columns colored green, blue, green, and yellow. V^T has rows colored green, blue, and yellow. The result is a 3x3 matrix where each element is the dot product of a row from V and a column from V^T .

4	0	3	1
3	0	1	2
4	5	0	3

\bullet

4	3	4
0	0	5
3	1	0

$=$

26	17	19
17	14	18
19	18	50

Gram matrix

Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

4	0	3	1
3	0	1	2
4	5	0	3

V

4	3	4
0	0	5
3	1	0

V^T

26	17	19
17	14	18
19	18	50

Gram matrix

Inner Product

Algebra

$$\vec{x} \cdot \vec{y} = \sum_1^n x_i y_i$$

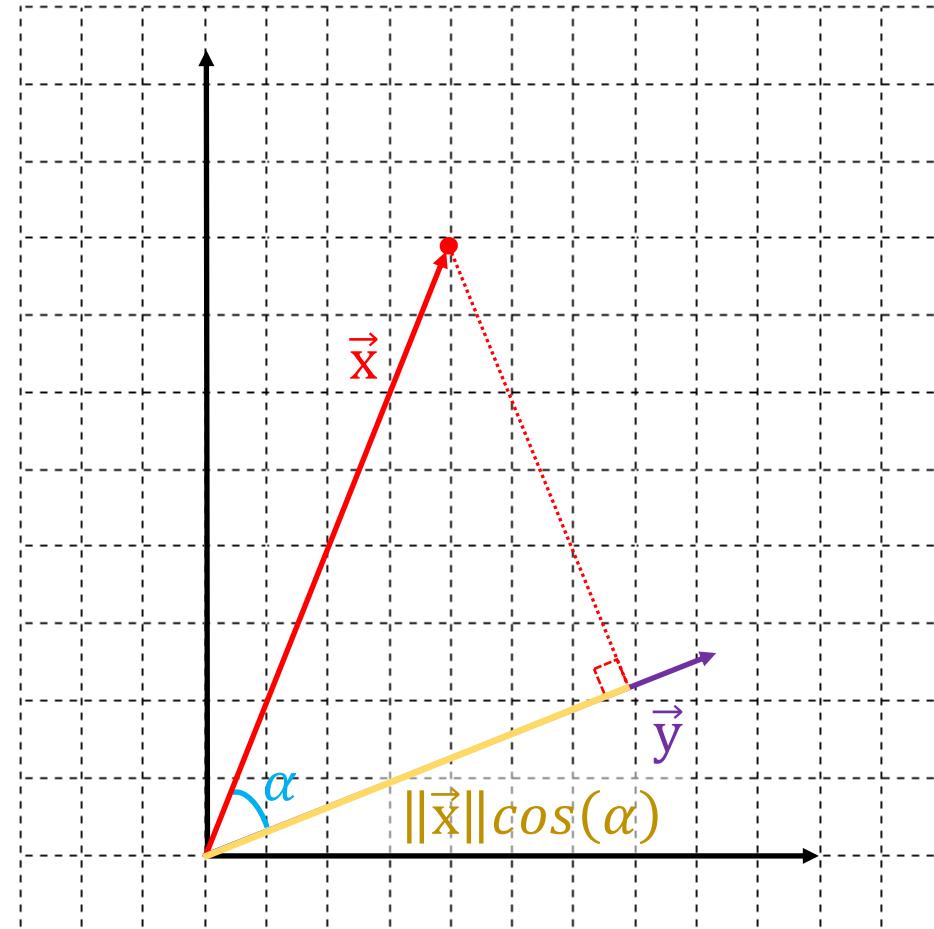
$$\begin{bmatrix} 4 & 0 & 3 & 1 \end{bmatrix} \cdot \begin{bmatrix} 4 \\ 0 \\ 3 \\ 1 \end{bmatrix} = 26$$

Geometry

$$\vec{x} \cdot \vec{y} = \|\vec{x}\| \|\vec{y}\| \cos(\alpha)$$

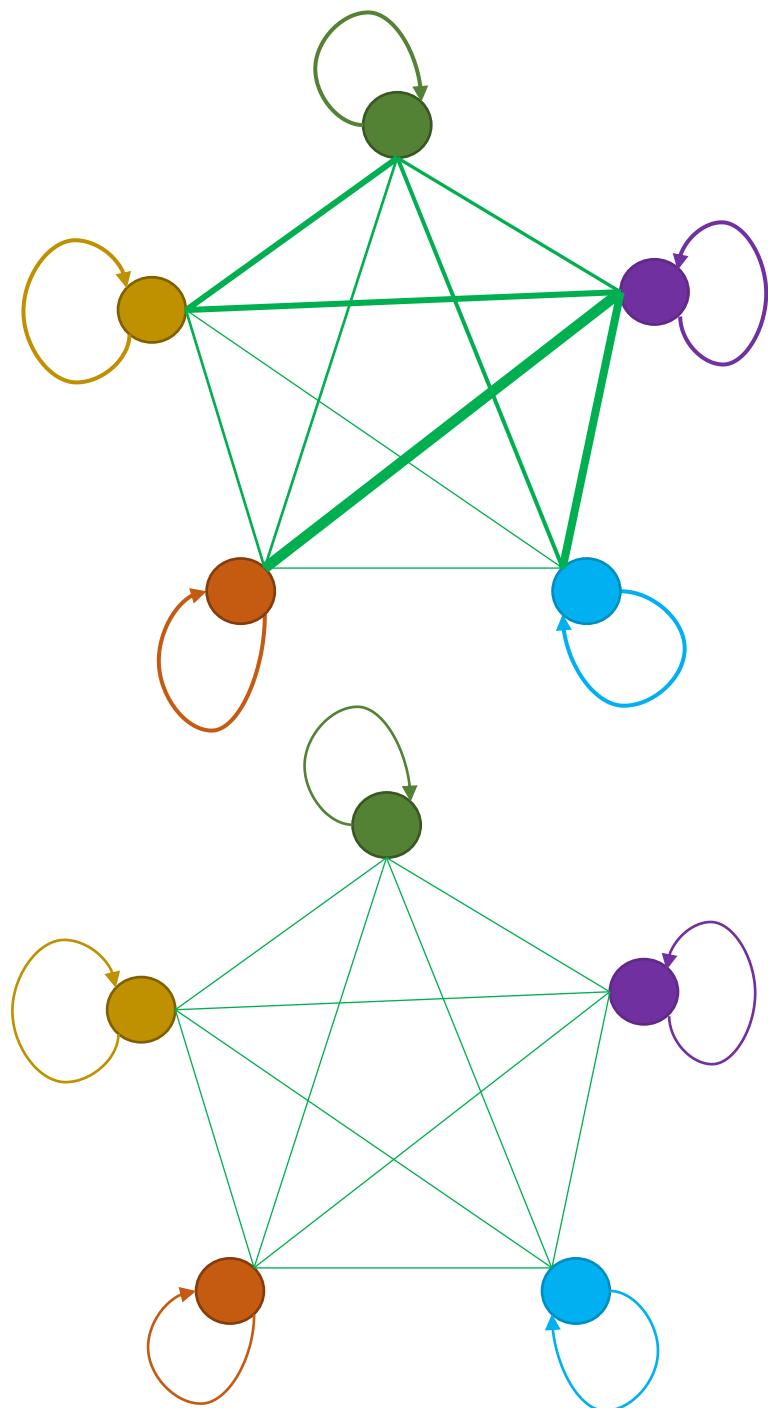
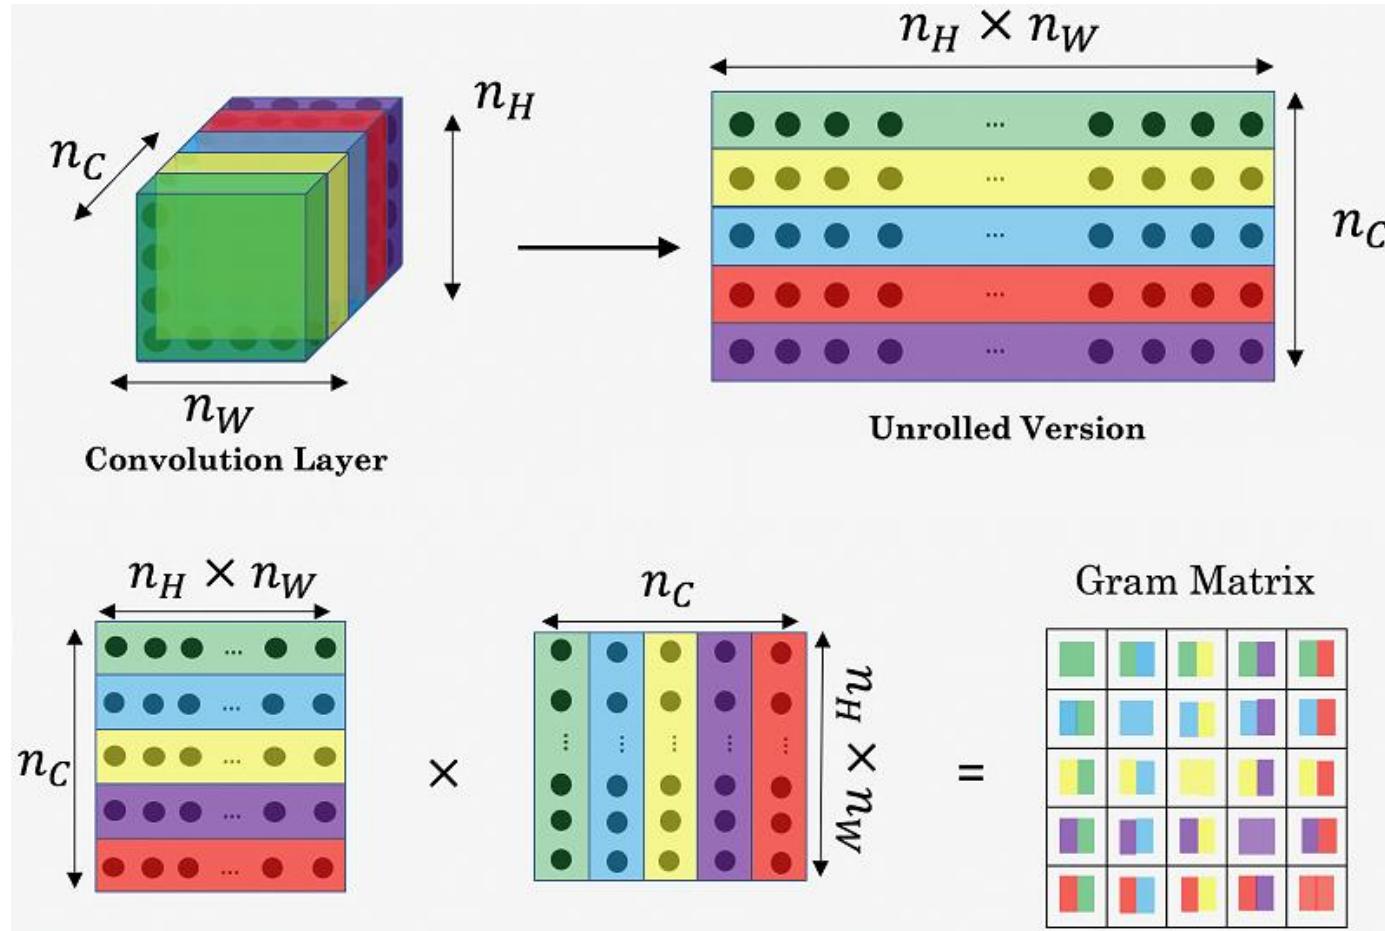
$$\begin{bmatrix} 3 & 0 & 4 & 1 \end{bmatrix} \cdot \begin{bmatrix} 3 \\ 0 \\ 4 \\ 1 \end{bmatrix} = 26$$

Position-invariant measure



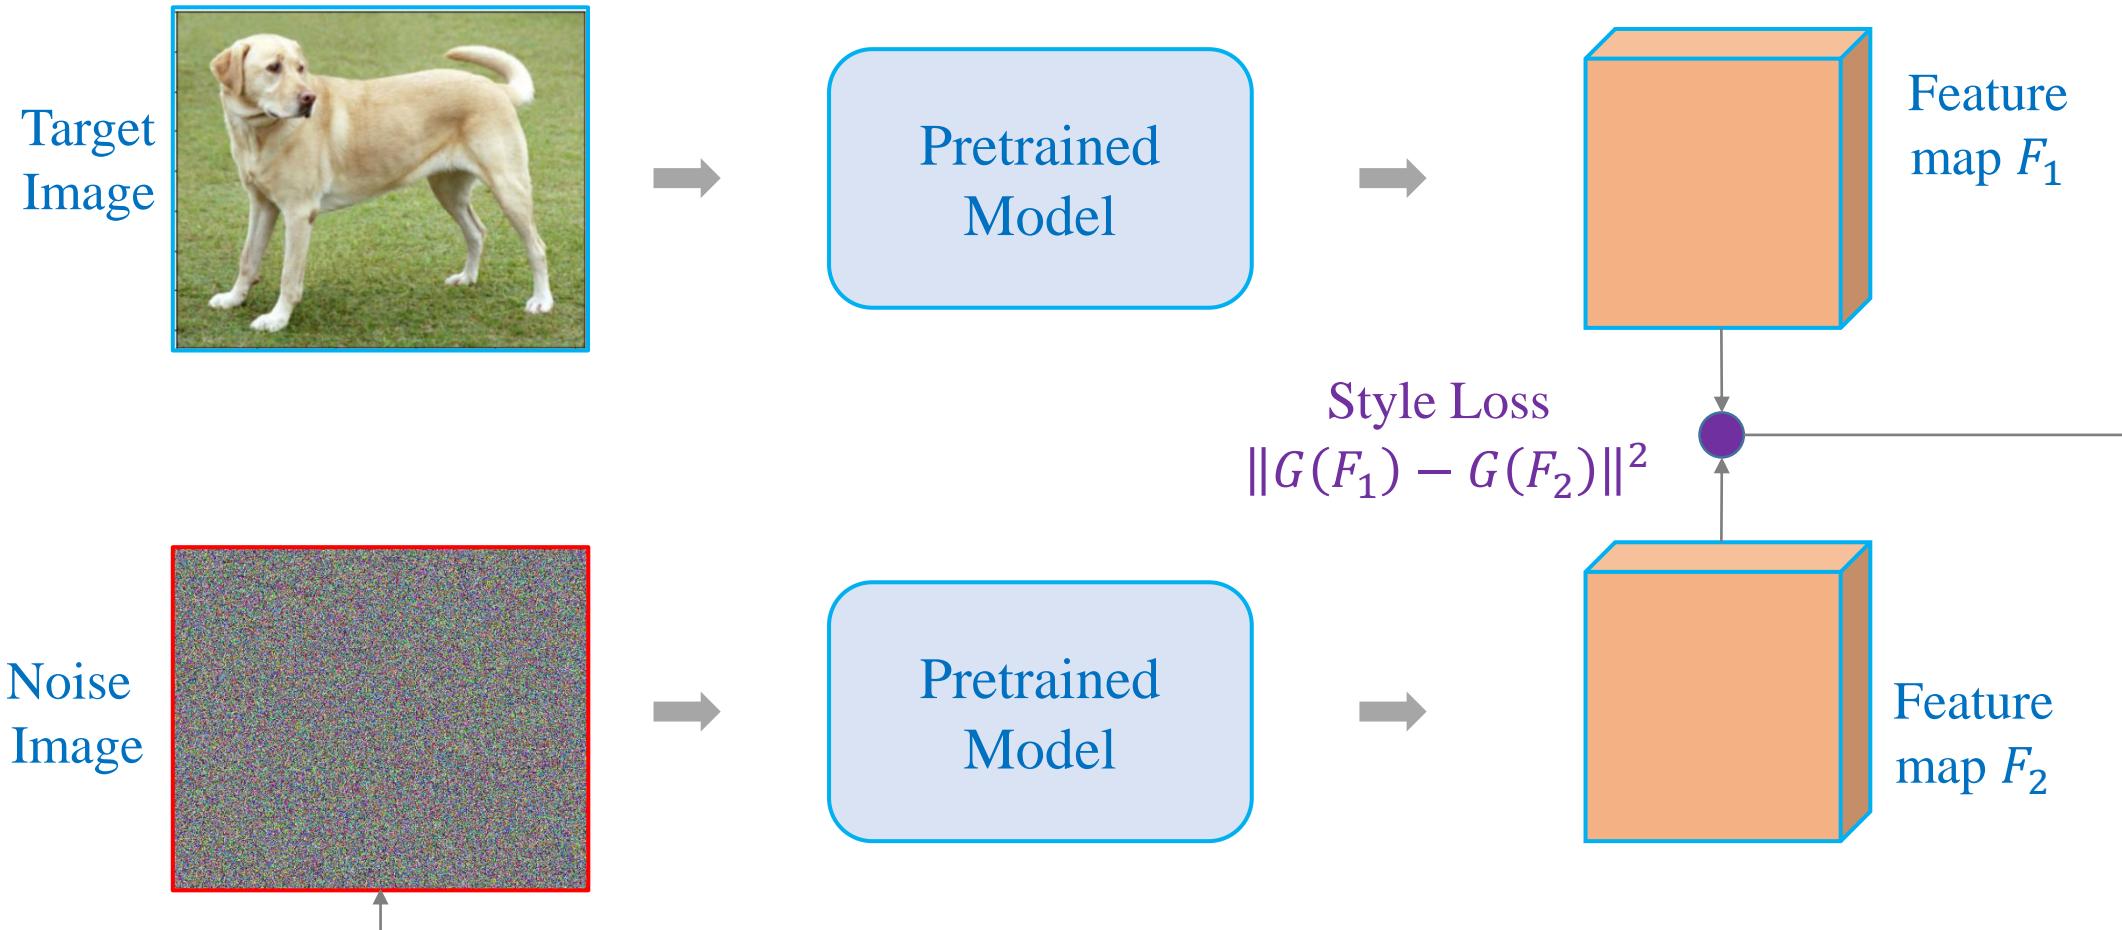
Review

❖ Gram matrix



Review

❖ Style Loss

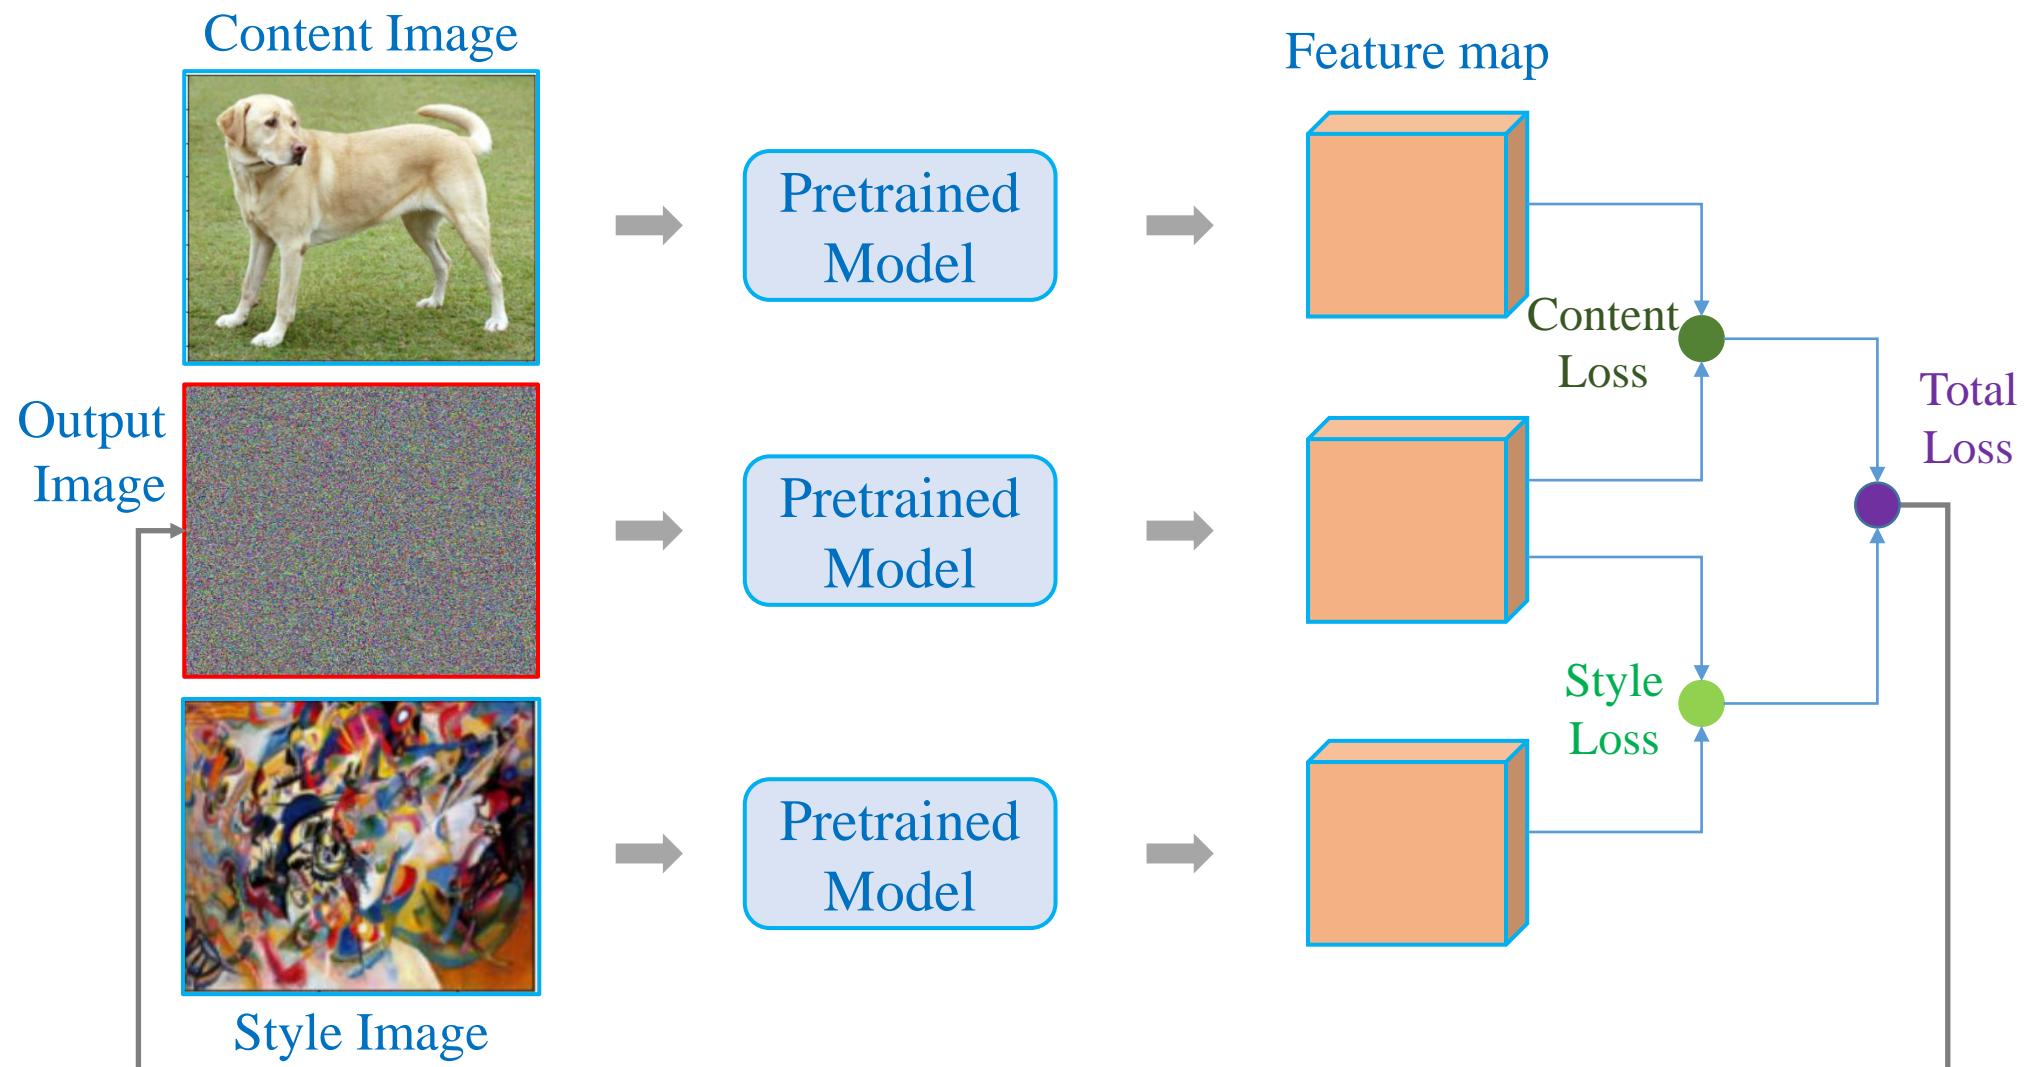


Outline

- Quick Review
- Content Loss+Style Loss
- Implementation
- Einstein Summation
- Extensions

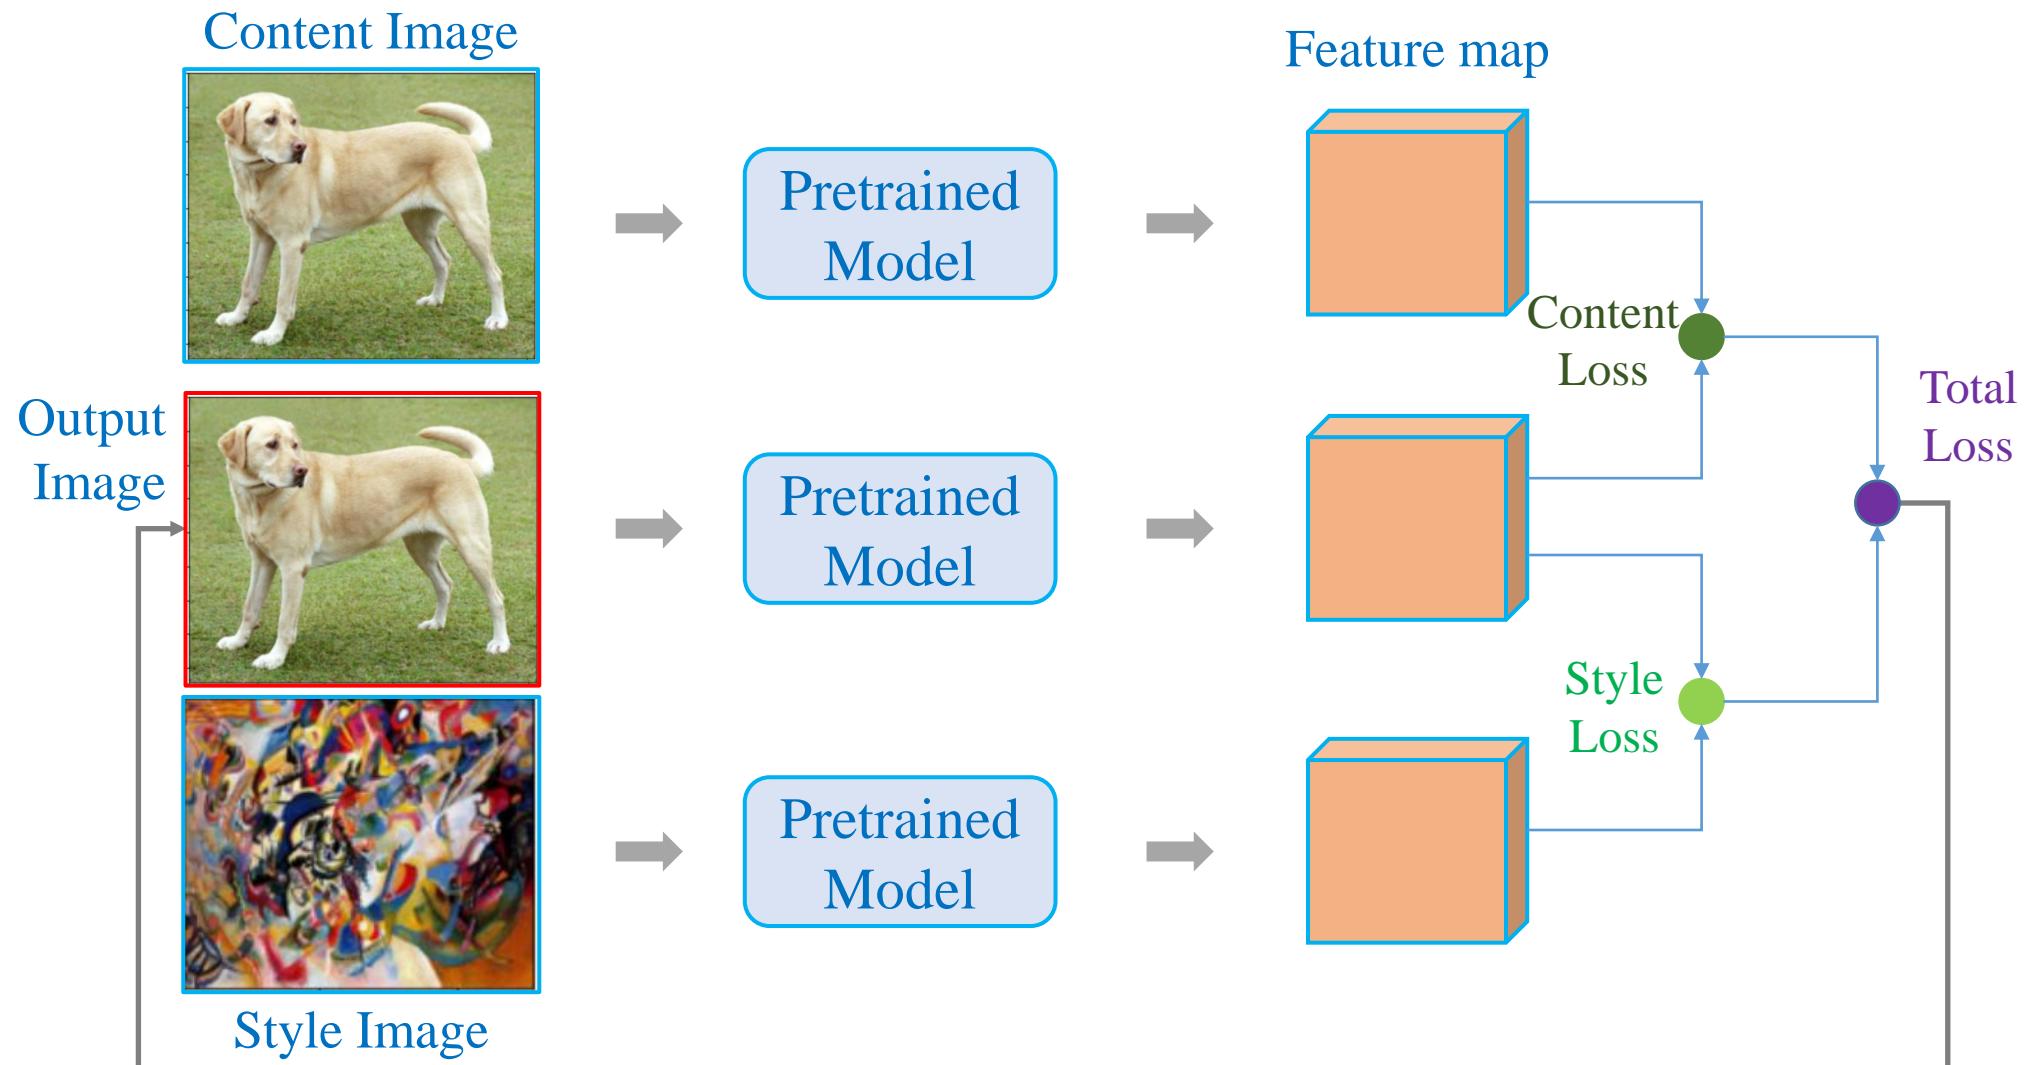
Style Transfer

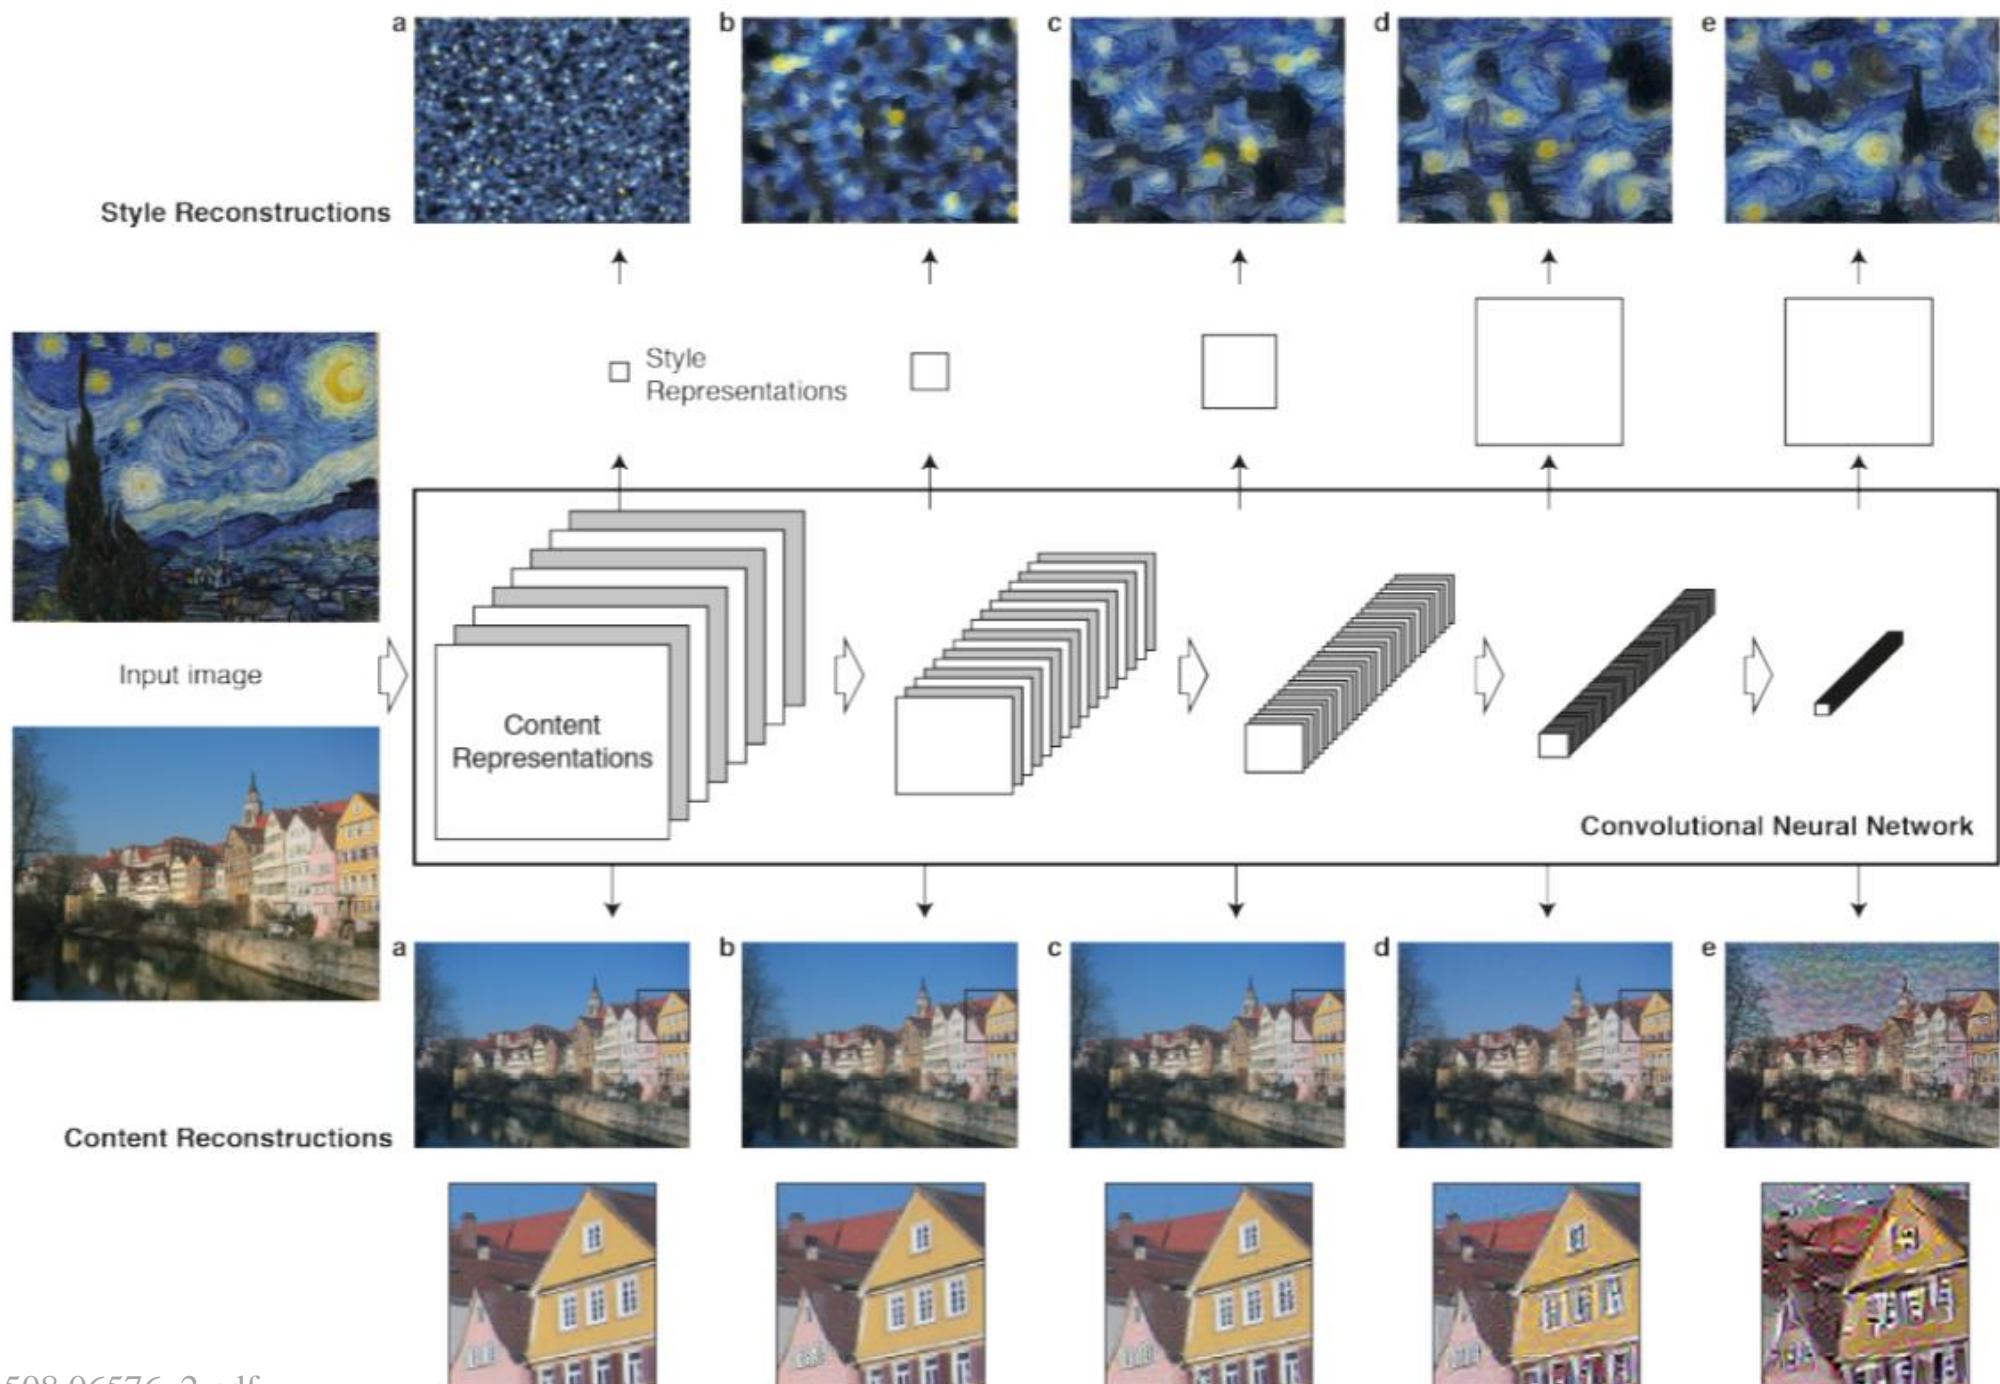
❖ Content Loss + Style Loss



Style Transfer

- ❖ Using the content image as initialization for the output image

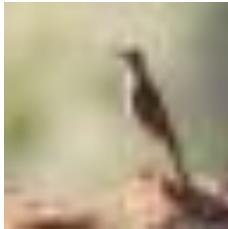




Style Transfer

❖ Example

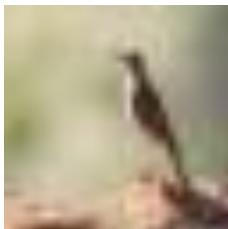
Content image



Shape=(3, 32, 32)

212	215	...	128	120	204	207	...	131	123	168	171	...	112	104
214	217	...	123	116	206	209	...	126	119	170	172	...	107	100
...
223	214	...	66	79	182	173	...	45	55	150	141	...	44	55
232	195	...	102	94	189	152	...	82	72	155	118	...	83	74

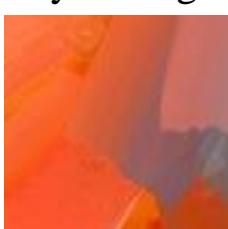
Output image



Shape=(3, 32, 32)

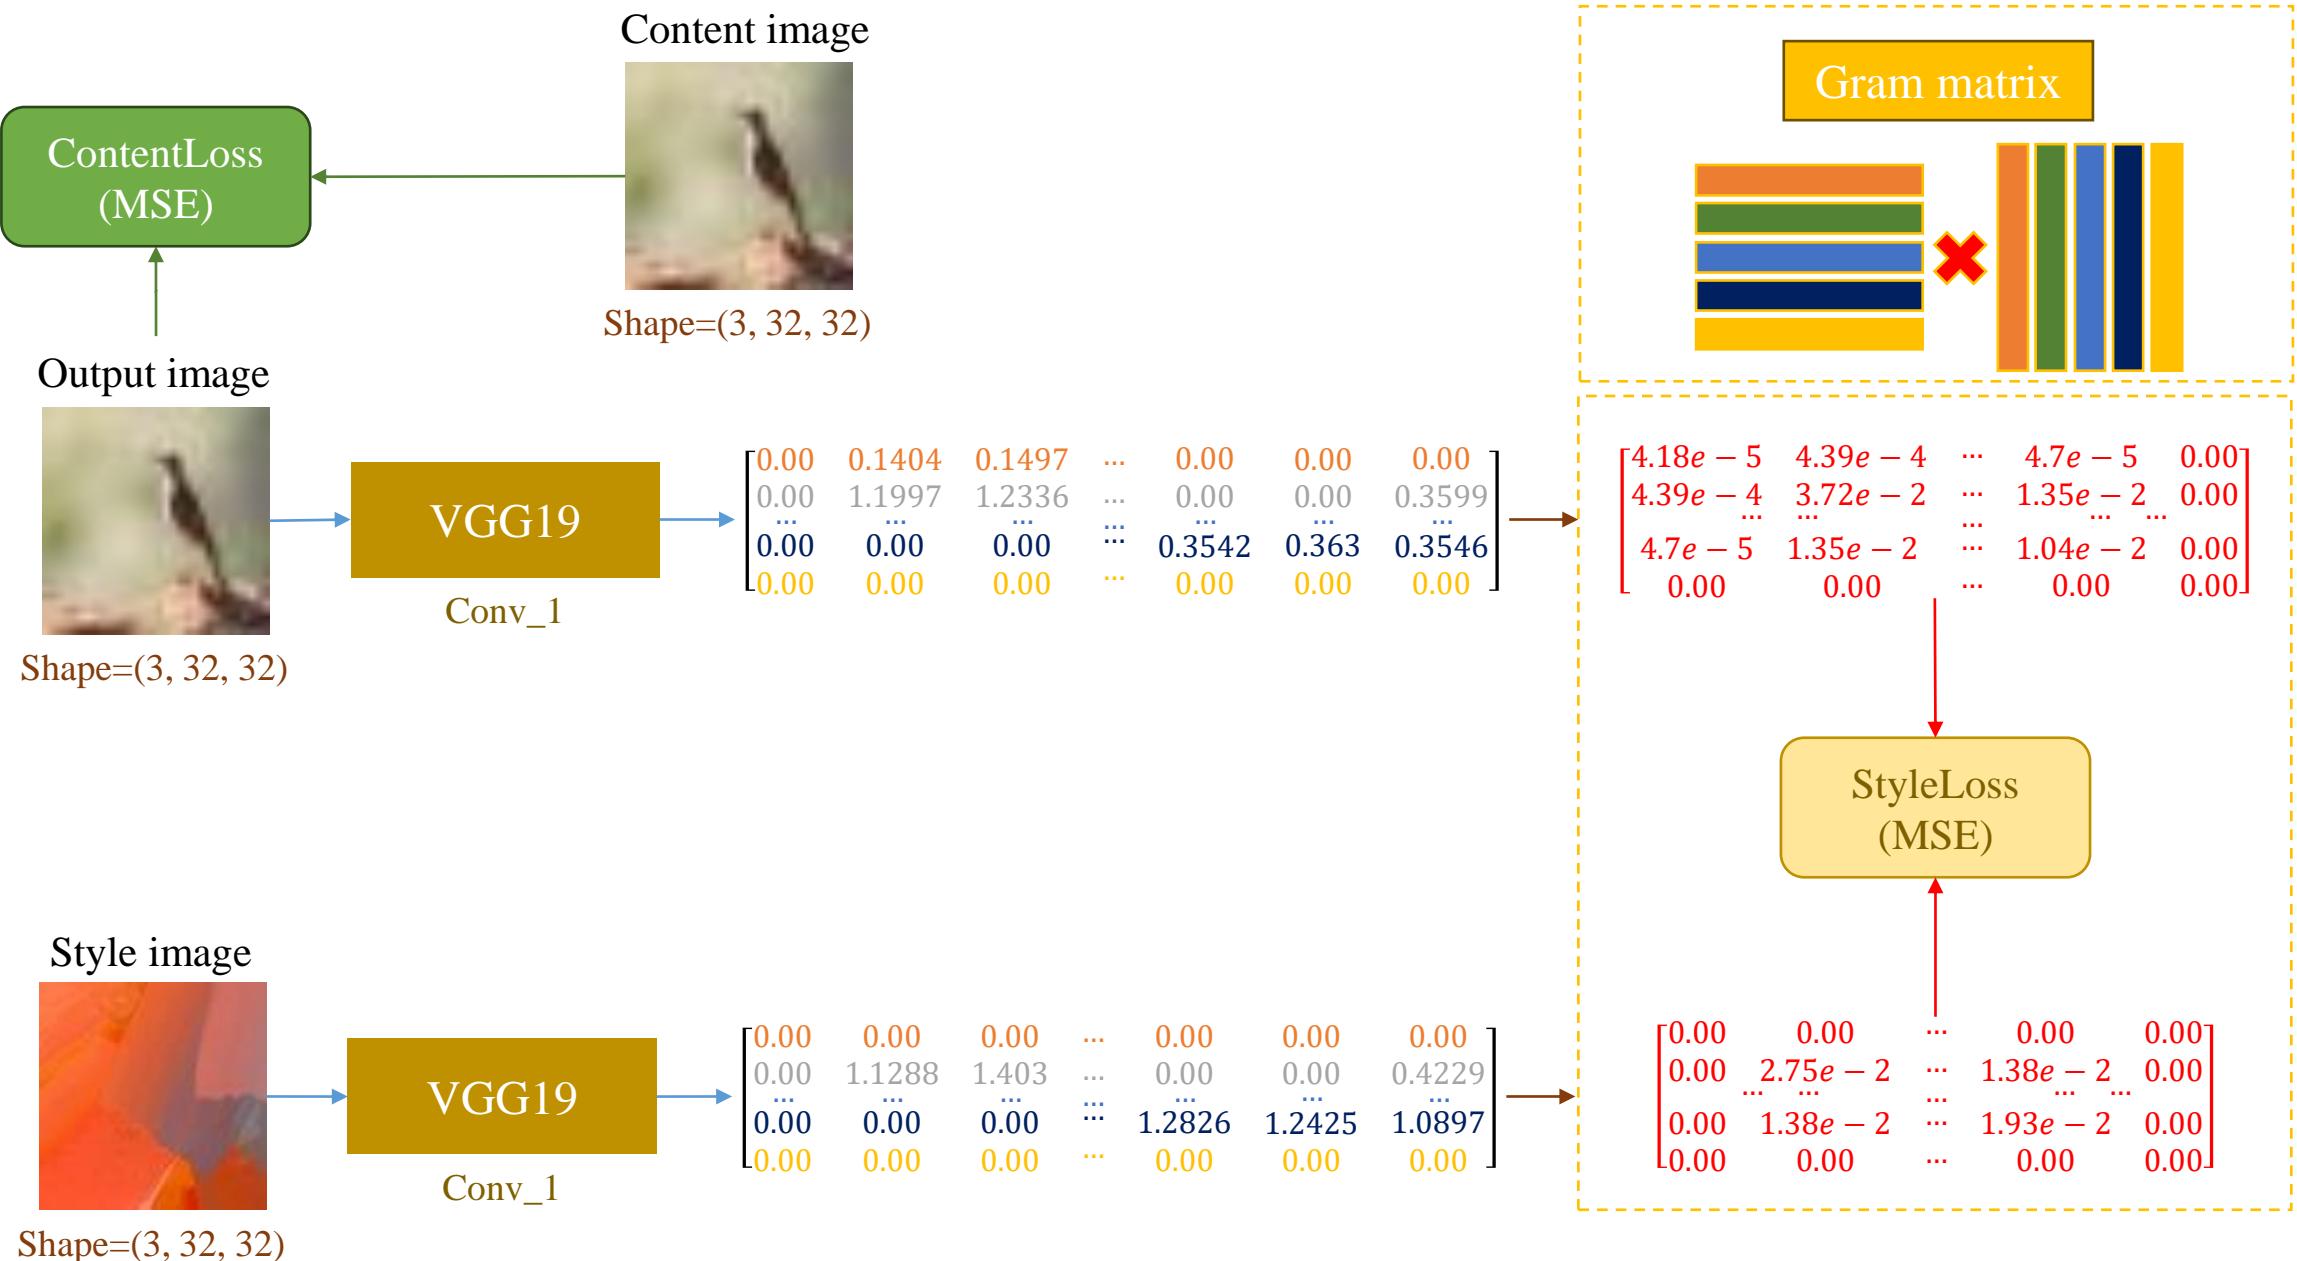
212	215	...	128	120	204	207	...	131	123	168	171	...	112	104
214	217	...	123	116	206	209	...	126	119	170	172	...	107	100
...
223	214	...	66	79	182	173	...	45	55	150	141	...	44	55
232	195	...	102	94	189	152	...	82	72	155	118	...	83	74

Style image



Shape=(3, 32, 32)

251	253	...	197	193	123	124	...	111	113	72	75	...	102	106
253	254	...	197	192	121	123	...	111	113	72	75	...	104	107
...
252	254	...	183	181	85	87	...	65	66	43	45	...	24	28
253	253	...	184	182	86	86	...	64	65	44	44	...	22	26



ContentLoss
(MSE)



ContentWeight
(1)

StyleLoss
(MSE)

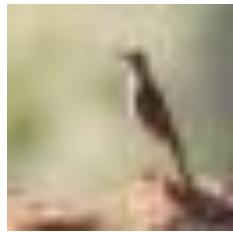


StyleWeight
(1e6)

TotalLoss (L)

Note: the pixel value of the output image when training will be normalized, not in uint8 type.

Output image



Shape=(3, 32, 32)

$$\begin{bmatrix} 212 & 215 & \cdots & 128 & 120 \\ 214 & 217 & \cdots & 123 & 116 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 223 & 214 & \cdots & 66 & 79 \\ 232 & 195 & \cdots & 102 & 94 \end{bmatrix} \xrightarrow{-0.05 \times} \begin{bmatrix} 0.0231 & -0.0083 & \cdots & -0.037 & -0.1309 \\ 0.0434 & 0.0678 & \cdots & 0.0996 & -0.0534 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ -0.017 & -0.0192 & \cdots & 0.0358 & -0.0135 \\ 0.0175 & 0.0347 & \cdots & 0.1084 & 0.0578 \end{bmatrix}$$

$$\begin{bmatrix} 204 & 207 & \cdots & 131 & 123 \\ 206 & 209 & \cdots & 126 & 119 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 182 & 173 & \cdots & 45 & 55 \\ 189 & 152 & \cdots & 82 & 72 \end{bmatrix} \xrightarrow{-0.05 \times} \begin{bmatrix} 0.0407 & 0.0107 & \cdots & -0.0027 & -0.1198 \\ 0.0652 & 0.01011 & \cdots & 0.1118 & -0.0578 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 0.022 & 0.0262 & \cdots & 0.0299 & -0.0281 \\ 0.0164 & 0.0203 & \cdots & 0.0662 & 0.0242 \end{bmatrix}$$

$$\begin{bmatrix} 168 & 171 & \cdots & 112 & 104 \\ 170 & 172 & \cdots & 107 & 100 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 150 & 141 & \cdots & 44 & 55 \\ 155 & 118 & \cdots & 83 & 74 \end{bmatrix} \xrightarrow{-0.05 \times} \begin{bmatrix} 0.0245 & 0.0201 & \cdots & 0.0312 & -0.0136 \\ 0.0283 & 0.0521 & \cdots & 0.04 & -0.003 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 0.0382 & 0.0378 & \cdots & 0.0021 & -0.0126 \\ -0.0006 & -0.009 & \cdots & -0.021 & -0.0168 \end{bmatrix}$$



$$\begin{bmatrix} 199 & 227 & \cdots & 140 & 132 \\ 201 & 204 & \cdots & 110 & 128 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 235 & 226 & \cdots & 53 & 91 \\ 219 & 182 & \cdots & 89 & 81 \end{bmatrix}$$

$$\begin{bmatrix} 191 & 194 & \cdots & 143 & 135 \\ 193 & 196 & \cdots & 113 & 131 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 169 & 160 & \cdots & 32 & 67 \\ 176 & 139 & \cdots & 69 & 59 \end{bmatrix}$$

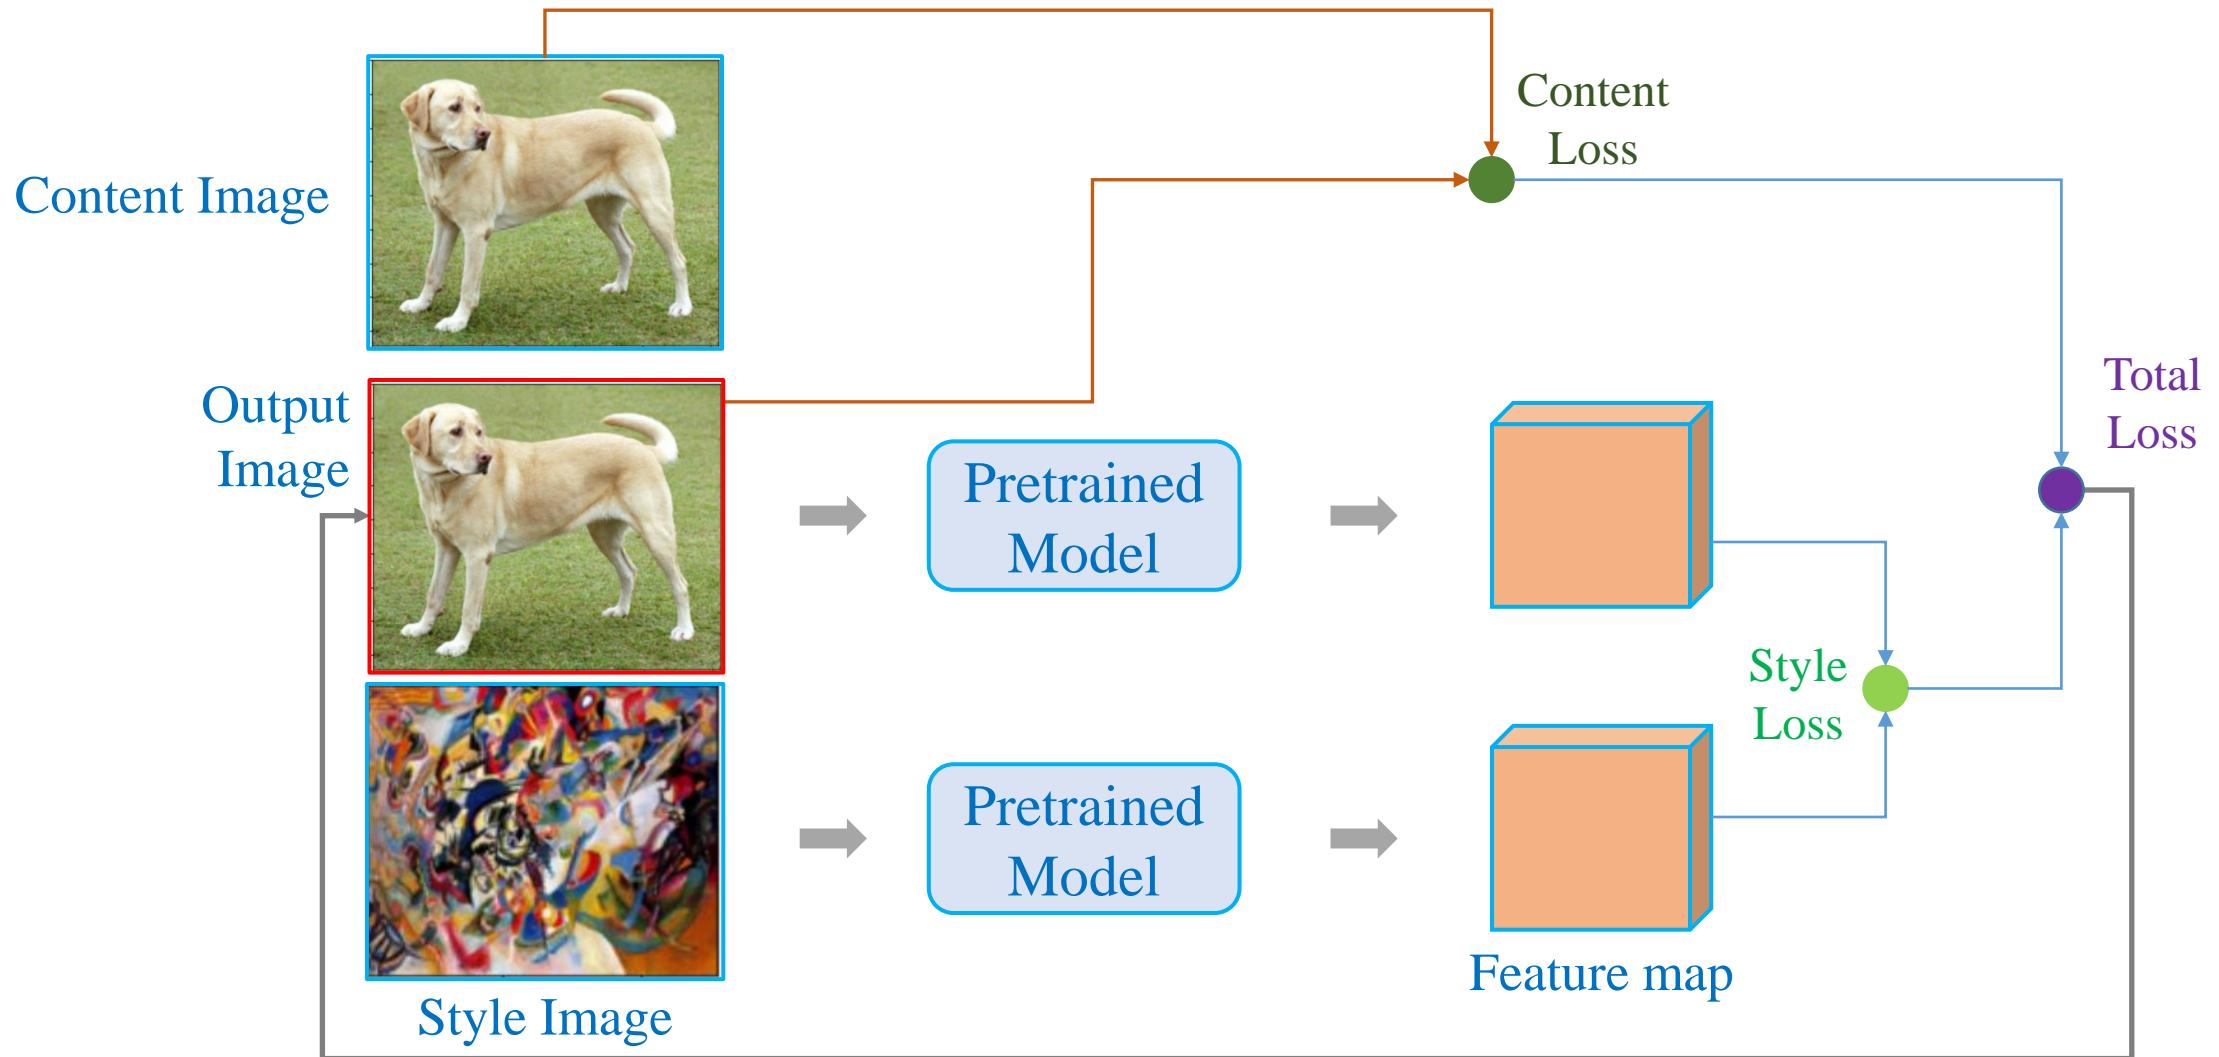
$$\begin{bmatrix} 155 & 158 & \cdots & 99 & 116 \\ 157 & 159 & \cdots & 94 & 112 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ 137 & 128 & \cdots & 31 & 67 \\ 167 & 130 & \cdots & 95 & 86 \end{bmatrix}$$

Outline

- Quick Review
- Content Loss+Style Loss
- Implementation
- Einstein Summation
- Extensions

Style Transfer

❖ Content Loss + Style Loss



Style Transfer

(1) Data Preparation

Content Image



Output
Image



Style Image

```
from PIL import Image
import torchvision.transforms as transforms

imsize = 256
img_transforms = transforms.Compose([
    transforms.Resize((imsize, imszie)),
    transforms.ToTensor()
])

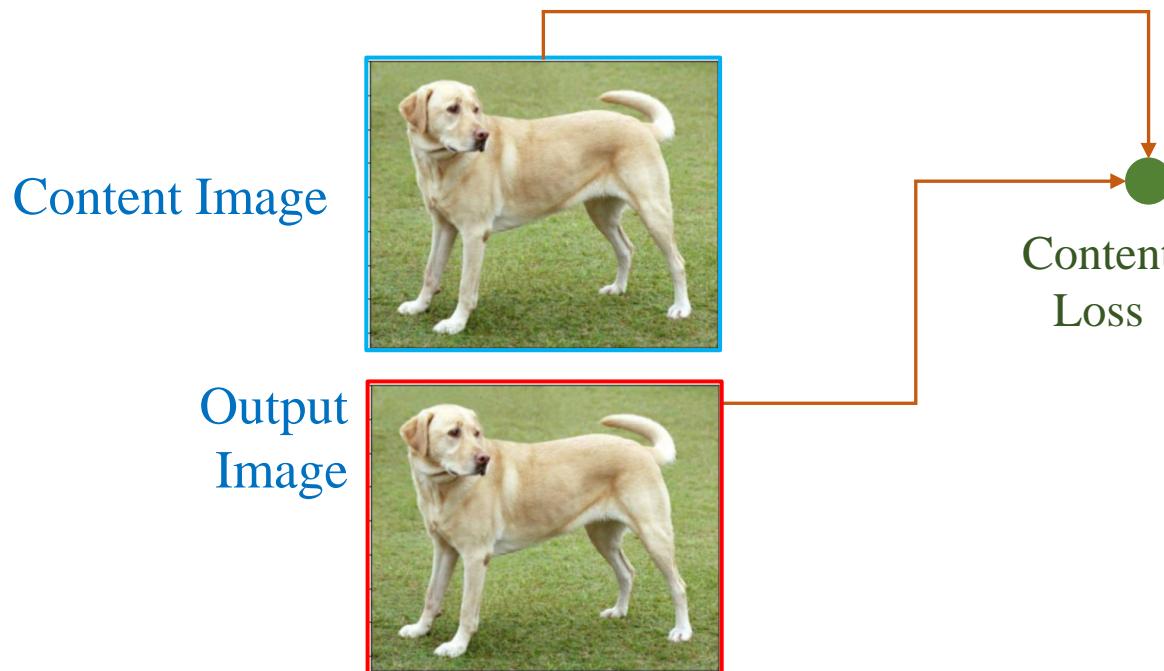
def image_loader(image_name):
    image = Image.open(image_name)
    image = img_transforms(image).unsqueeze(0)
    return image.to(device, torch.float)

style_img = image_loader("style_img.jpg")
content_img = image_loader("content_img.jpg")

output_img = content_img.clone().requires_grad_(True).to(device)
```

Style Transfer

(2) Content Loss



```
# Content Loss
ContentLoss = nn.MSELoss()

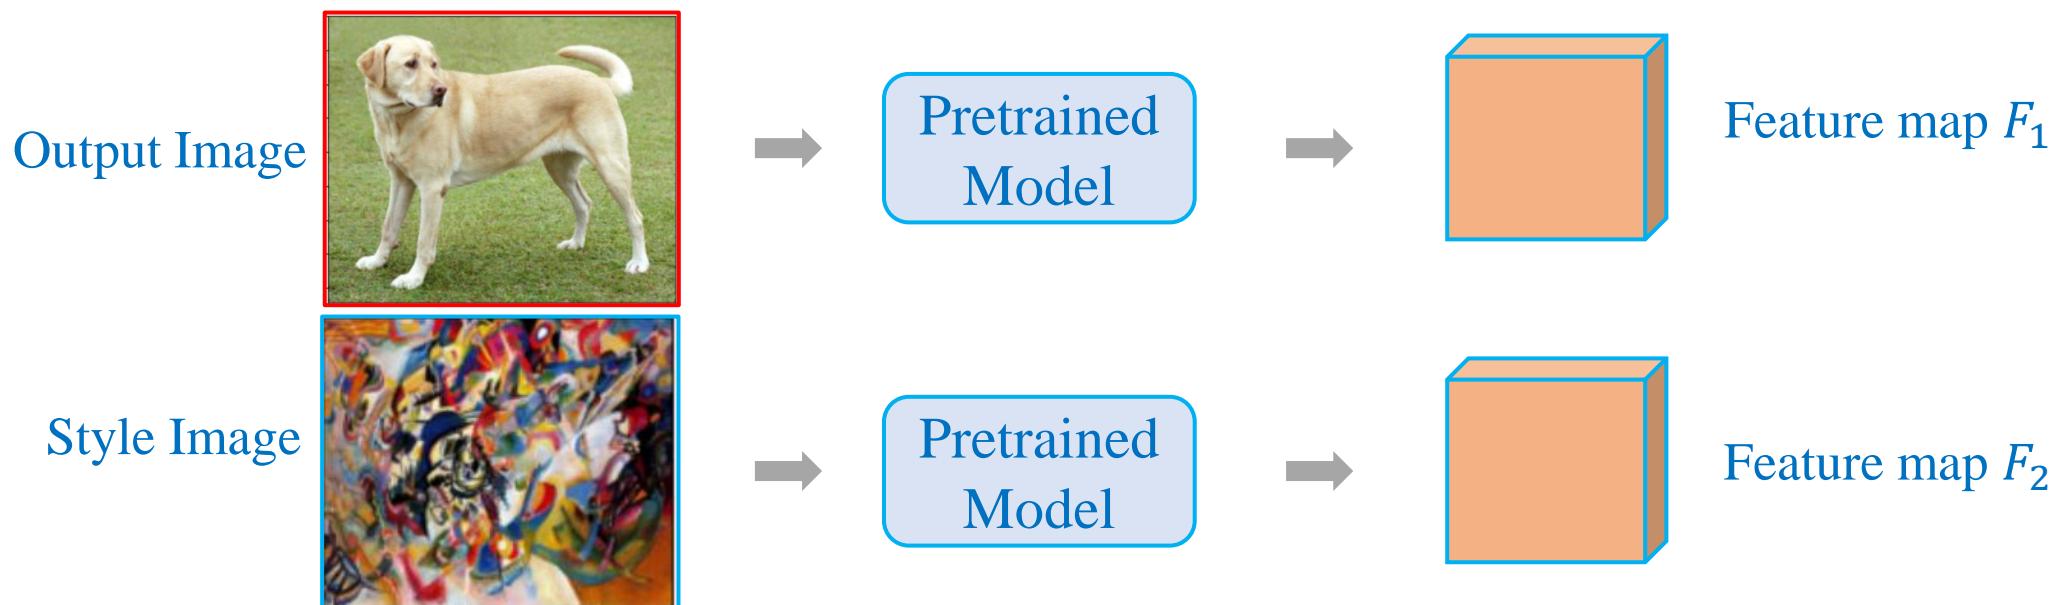
STEPS = 2000
for step in range(STEPS):
    optimizer.zero_grad()
    content_loss = ContentLoss(content_img,
                                output_img)

    # ...
    total_loss.backward()
    optimizer.step()
```

Style Transfer

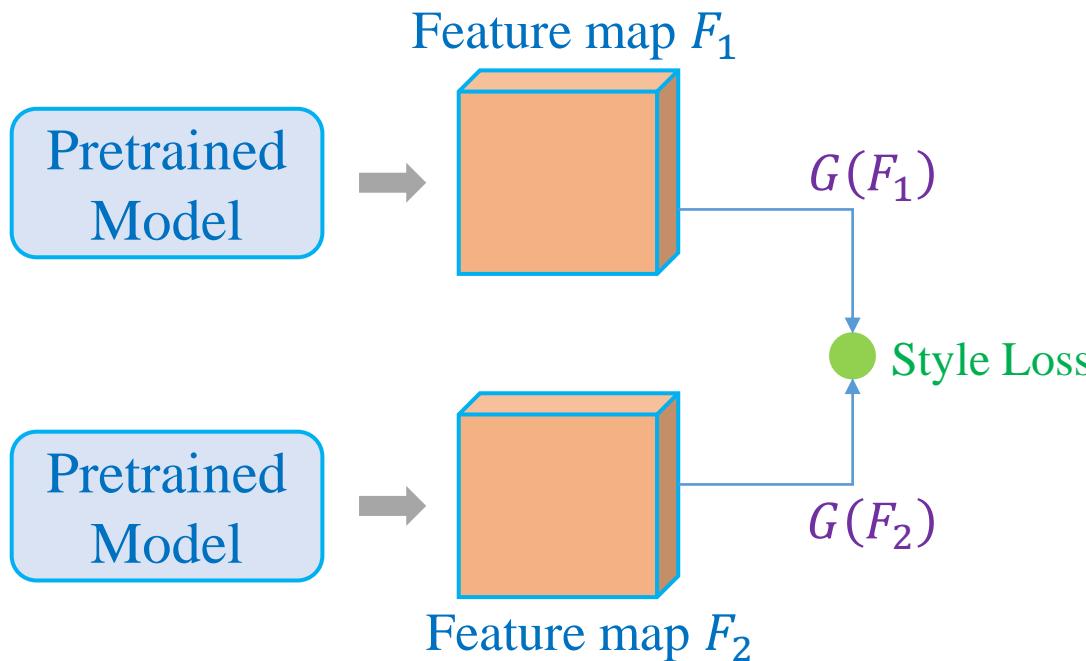
❖ Content Loss + Style Loss

```
VGG19_pretrained = vgg19(weights=VGG19_Weights.DEFAULT).features.eval()  
for param in VGG19_pretrained.parameters():  
    param.requires_grad_(False)  
  
# get outputs from the conv1 of the block 3  
model = VGG19_pretrained[:11]  
model.to(device)  
  
style_features = model(style_img)  
output_features = model(output_img)
```



Style Transfer

❖ Content Loss + Style Loss



```
def gram_matrix(tensor):
    a, b, c, d = tensor.size()
    tensor = tensor.view(a * b, c * d)
    G = torch.mm(tensor, tensor.t())
    return G.div(a * b * c * d)

# Style Loss
StyleLoss = nn.MSELoss()

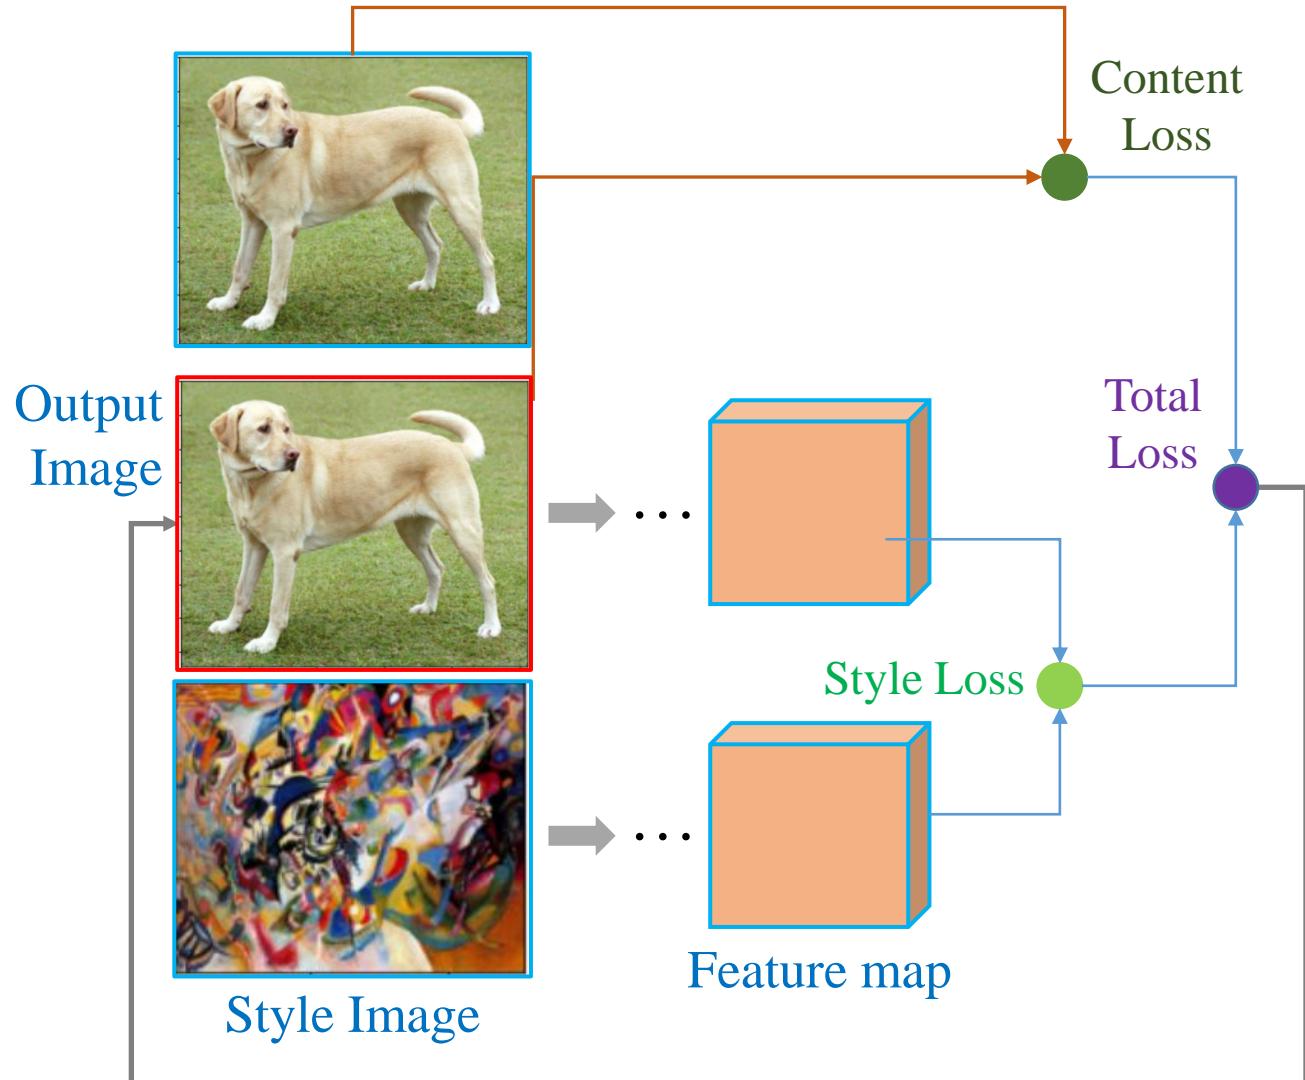
# compute Gram for the style image
style_features = model(style_img)
style_gram = gram_matrix(style_features)

STEPS = 2000
for step in range(STEPS):
    output_features = model(output_img)
    output_gram = gram_matrix(output_features)
    style_loss = StyleLoss(style_gram,
                          output_gram)

    # ...
```

Style Transfer

❖ Content Loss + Style Loss



```
style_features = model(style_img)
style_gram = gram_matrix(style_features)

STEPS = 2000
for step in range(STEPS):
    optimizer.zero_grad()

    # content loss
    content_loss = ContentLoss(content_img,
                                output_img)

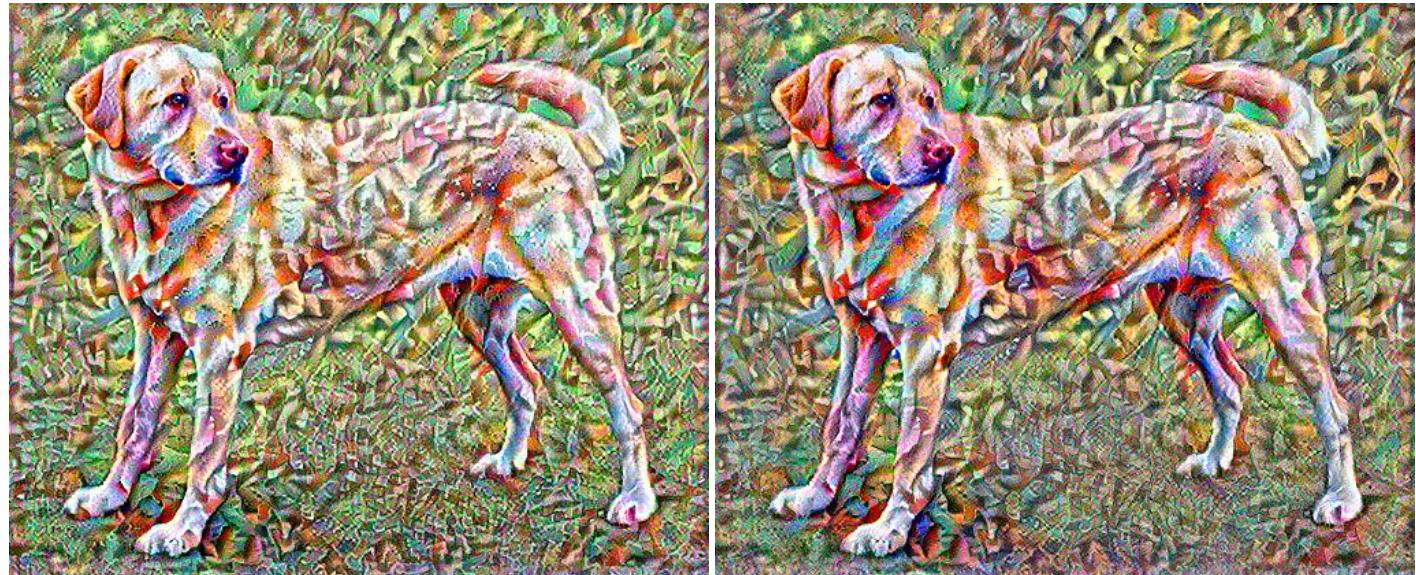
    # style loss
    output_features = model(output_img)
    output_gram = gram_matrix(output_features)
    style_loss = StyleLoss(style_gram,
                          output_gram)

    # total loss
    total_loss = content_loss + style_loss*weight
    total_loss.backward()
    optimizer.step()

    with torch.no_grad():
        output_img.clamp_(0, 1)
```

Style Transfer

❖ Content Loss + Style Loss



Outline

- Quick Review
- Content Loss+Style Loss
- Implementation
- Einstein Summation
- Extensions

Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

light green	light green	light green	light green
light blue	light blue	light blue	light blue
light yellow	light yellow	light yellow	light yellow



light green	light blue	light yellow
light green	light blue	light yellow
light green	light blue	light yellow



light orange	light blue	light orange
light orange	light blue	light orange
light orange	light blue	light orange

V

V^T

Gram matrix

```
import torch

x = torch.Tensor([[4, 0, 3, 1],
                  [3, 0, 1, 2],
                  [4, 5, 0, 3]])
G = torch.einsum('ij,jk->ik', x, x.t())
print(G)

tensor([[26., 17., 19.],
        [17., 14., 18.],
        [19., 18., 50.]])
```

4	0	3	1
3	0	1	2
4	5	0	3

V



4	3	4
0	0	5
3	1	0
1	2	3

V^T



26	17	19
17	14	18
19	18	50

Gram matrix

Einstein Summation

Dot product

$$\begin{matrix} 3 & 2 & 4 & 1 \end{matrix} \cdot \begin{matrix} 2 \\ 1 \\ 0 \\ 3 \end{matrix} = 11$$

u v

```
import torch

u = torch.Tensor([3, 2, 4, 1])
v = torch.Tensor([2, 1, 0, 3])

r = torch.einsum('i,i->', u, v)
print(r)

tensor(11.)
```

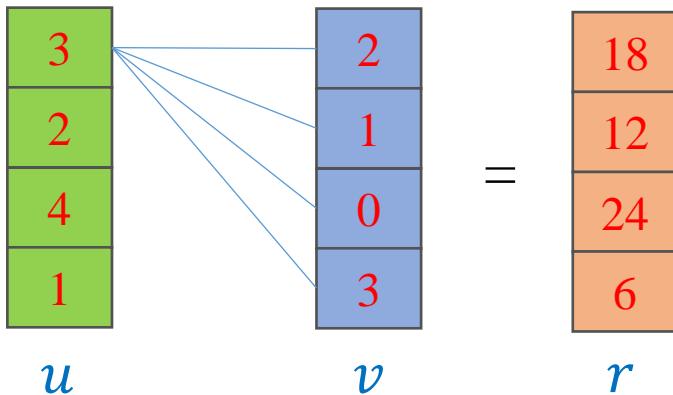
$i, i \rightarrow$

$$i \rightarrow \begin{matrix} 3 \\ 2 \\ 4 \\ 1 \end{matrix} \quad i \rightarrow \begin{matrix} 2 \\ 1 \\ 0 \\ 3 \end{matrix}$$

u v

$$r = \sum_i u_i v_i = u_1 v_1 + u_2 v_2 + \dots$$

Einstein Summation

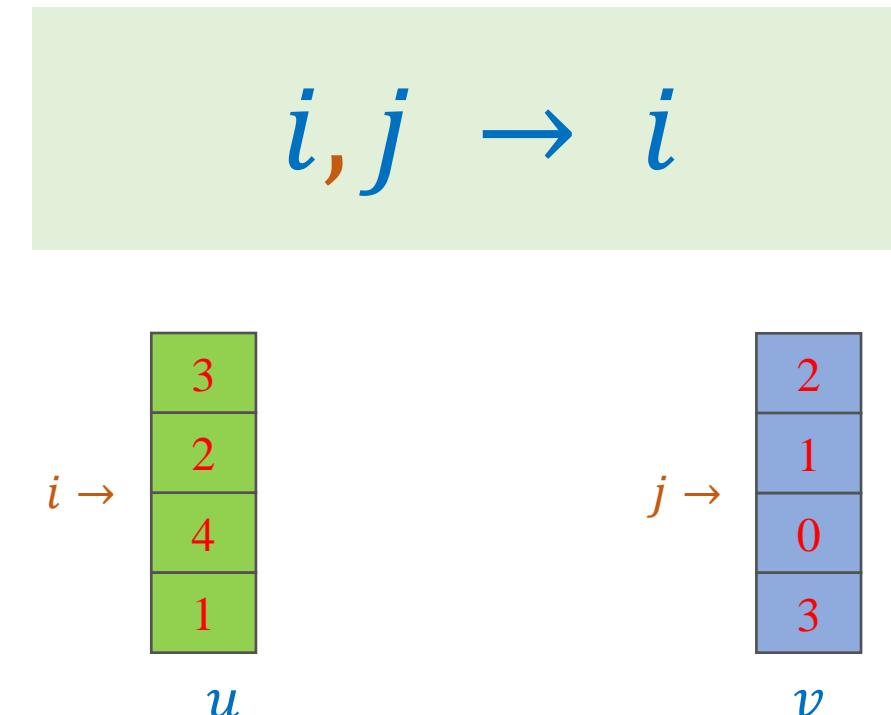


```
import torch

u = torch.Tensor([3, 2, 4, 1])
v = torch.Tensor([2, 1, 0, 3])

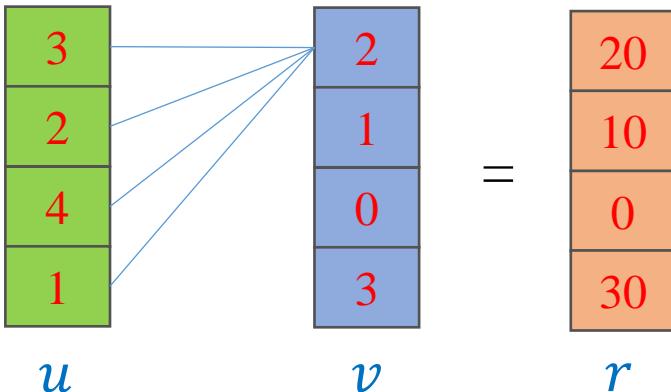
r = torch.einsum('i,j->i', u, v)
print(r)

tensor([18., 12., 24., 6.])
```



$$r_i = \sum_j u_i v_j = u_i v_1 + u_i v_2 + \dots$$

Einstein Summation



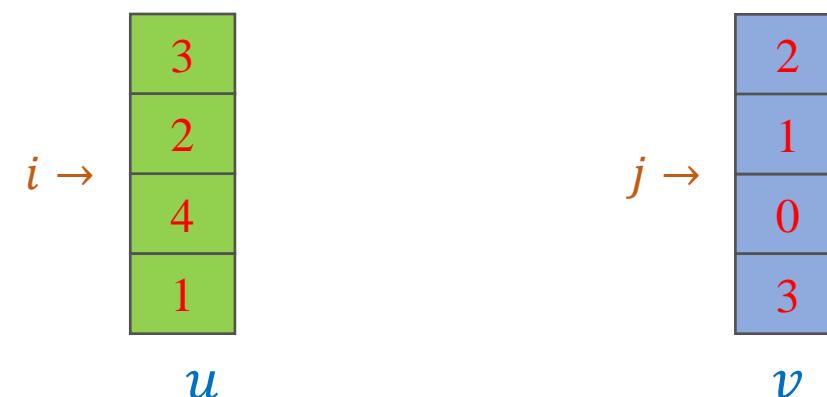
```
import torch

u = torch.Tensor([3, 2, 4, 1])
v = torch.Tensor([2, 1, 0, 3])

r = torch.einsum('i,j->j', u, v)
print(r)

tensor([20., 10., 0., 30.])
```

$$i, j \rightarrow j$$



$$r_j = \sum_i u_i v_j = u_1 v_j + u_1 v_j + \dots$$

Einstein Summation

Outer product

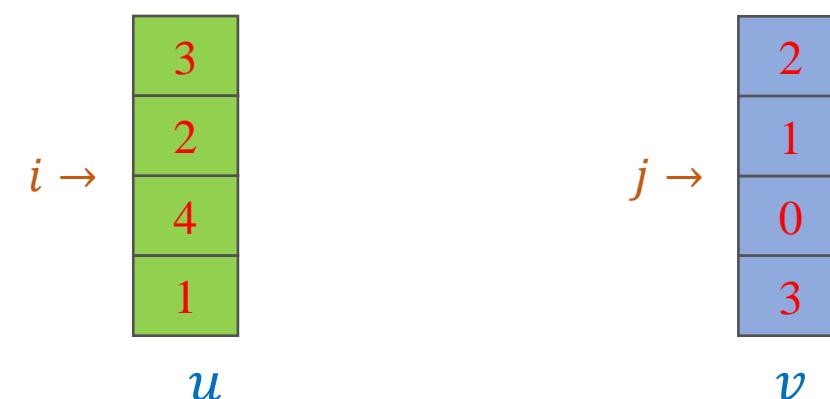
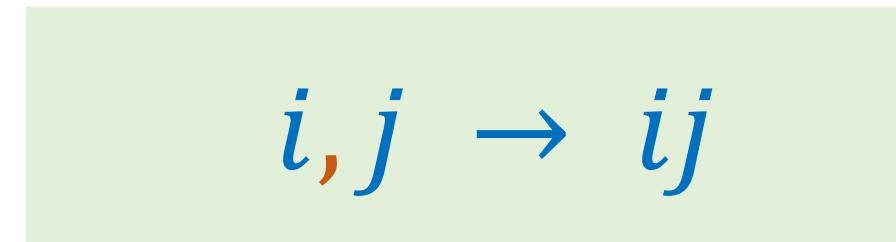
$$u \begin{pmatrix} 3 \\ 2 \\ 4 \\ 1 \end{pmatrix} \bullet v \begin{pmatrix} 2 & 1 & 0 & 3 \end{pmatrix} = r \begin{pmatrix} 6 & 3 & 0 & 9 \\ 4 & 2 & 0 & 6 \\ 8 & 4 & 0 & 12 \\ 2 & 1 & 0 & 3 \end{pmatrix}$$

```

u = torch.Tensor([3, 2, 4, 1])
v = torch.Tensor([2, 1, 0, 3])

r = torch.einsum('i,j->ij', u, v)
print(r)

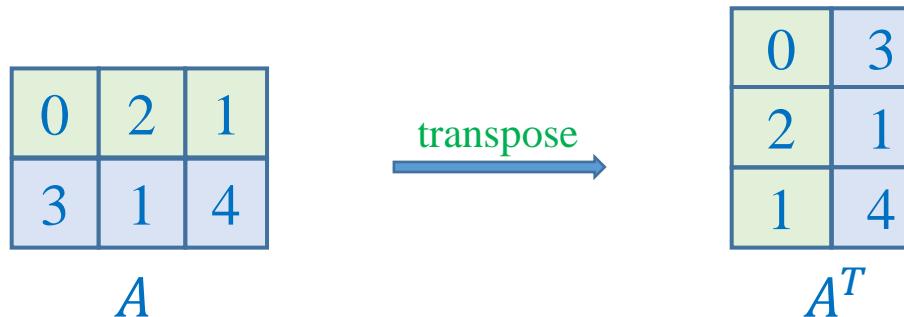
tensor([[ 6.,  3.,  0.,  9.],
        [ 4.,  2.,  0.,  6.],
        [ 8.,  4.,  0., 12.],
        [ 2.,  1.,  0.,  3.]])
    
```



$$r_{ij} = u_i v_j$$

Einstein Summation

Transpose

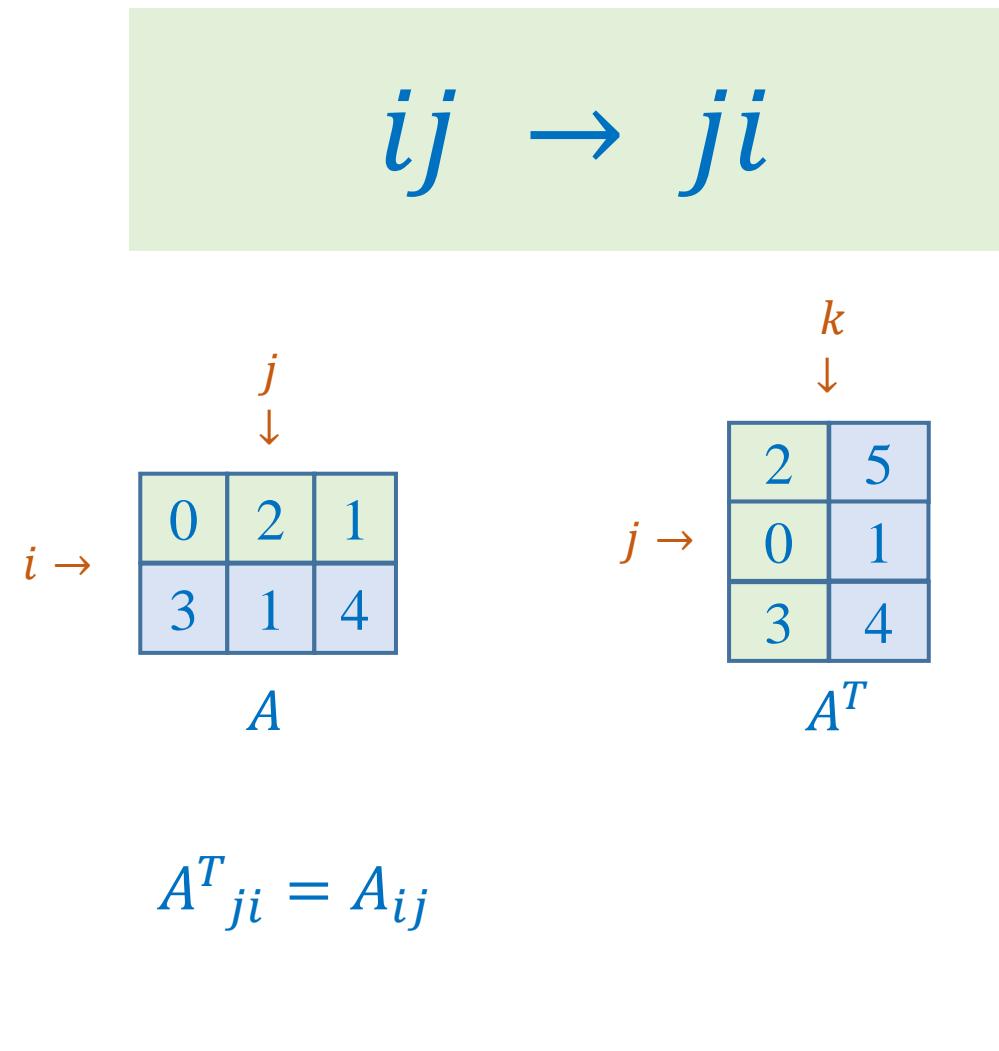


```
import torch

A = torch.Tensor([[0, 2, 1],
                  [3, 1, 4]])

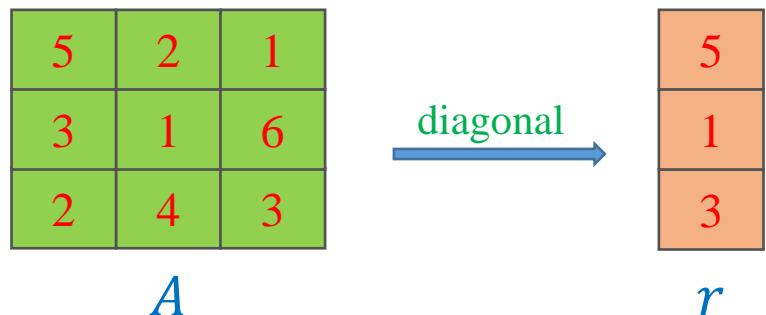
T = torch.einsum('ij->ji', A)
print(T)

tensor([[0., 3.],
        [2., 1.],
        [1., 4.]])
```



Einstein Summation

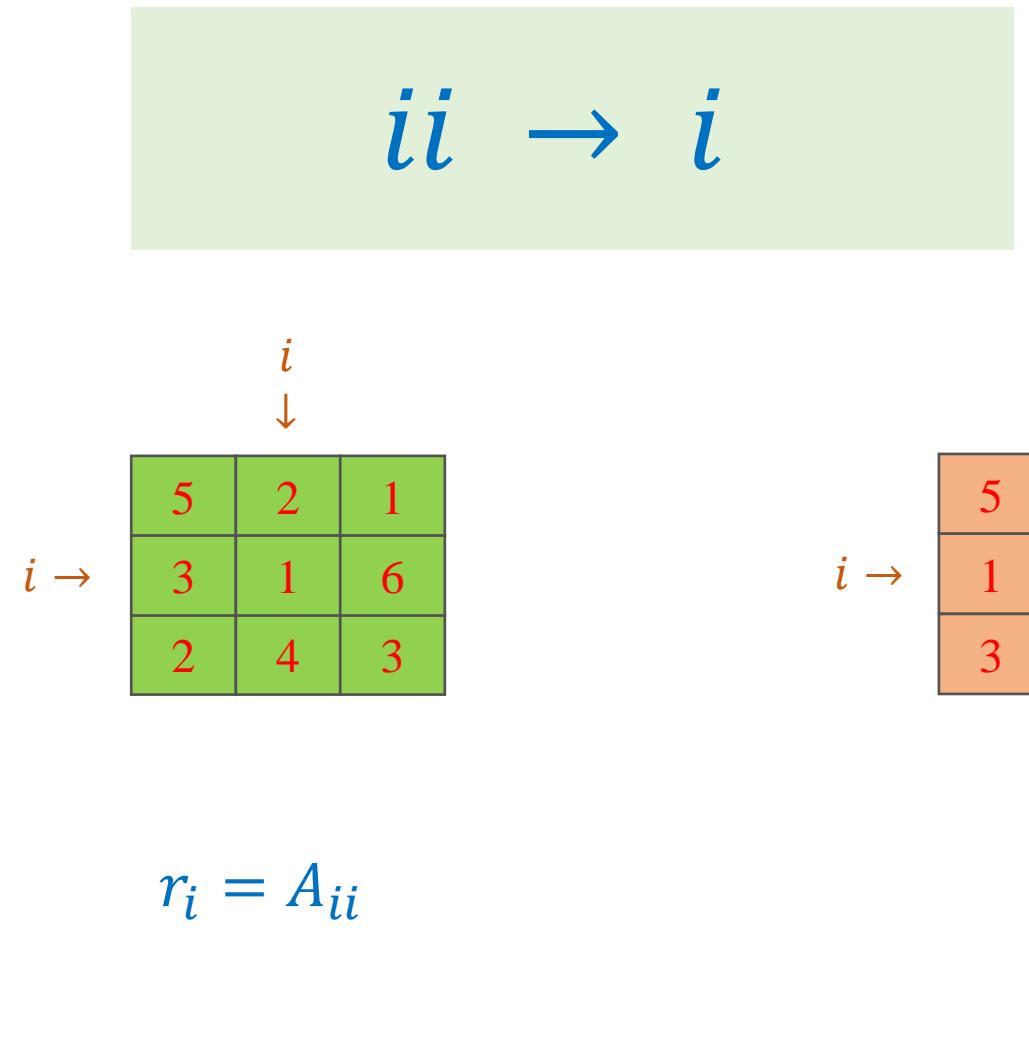
Diagonal



```
import torch

A = torch.Tensor([[5, 2, 1],
                  [3, 1, 6],
                  [2, 4, 3]])
r = torch.einsum('ii->i', A)
print(r)

tensor([5., 1., 3.])
```



Einstein Summation

Trace

5	2	1
3	1	6
2	4	3

A

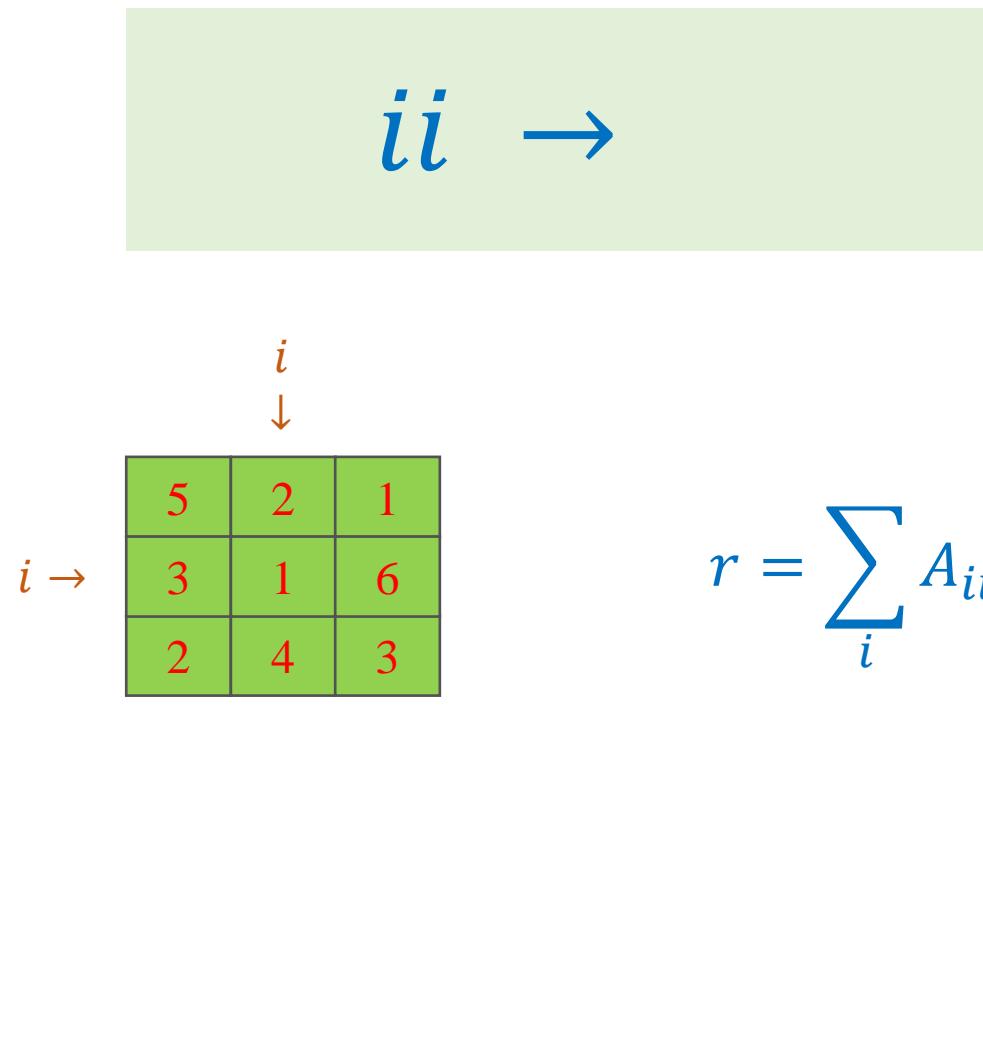
$$\xrightarrow{\text{trace}} \begin{matrix} 9 \\ r \end{matrix}$$

```
import torch

A = torch.Tensor([[5, 2, 1],
                  [3, 1, 6],
                  [2, 4, 3]])

r = torch.einsum('ii->', A)
print(r)

tensor(9.)
```



Einstein Summation

Element-wise multiplication

$$\begin{array}{|c|c|c|} \hline 3 & 0 & 1 \\ \hline 2 & 2 & 3 \\ \hline \end{array} \odot \begin{array}{|c|c|c|} \hline 1 & 0 & 2 \\ \hline 3 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 3 & 0 & 2 \\ \hline 6 & 4 & 3 \\ \hline \end{array}$$

A B C

```
import torch

A = torch.Tensor([[3, 0, 1],
                  [2, 2, 3]])
B = torch.Tensor([[1, 0, 2],
                  [3, 2, 1]])
r = torch.einsum('ij,ij->ij', A, B)
print(r)

tensor([[3., 0., 2.],
        [6., 4., 3.]])
```

$$ij, ij \rightarrow ij$$

$$i \rightarrow \begin{array}{|c|c|c|} \hline 3 & 0 & 1 \\ \hline 2 & 2 & 3 \\ \hline \end{array} \quad j \downarrow \quad A$$
$$i \rightarrow \begin{array}{|c|c|c|} \hline 1 & 0 & 2 \\ \hline 3 & 2 & 1 \\ \hline \end{array} \quad j \downarrow \quad B$$
$$i \rightarrow \begin{array}{|c|c|c|} \hline 3 & 0 & 2 \\ \hline 6 & 4 & 3 \\ \hline \end{array} \quad j \downarrow \quad C$$

$$C_{ij} = A_{ij}B_{ij}$$

Einstein Summation

Matrix multiplication

$$\begin{array}{|c|c|c|} \hline 0 & 2 & 1 \\ \hline 3 & 1 & 4 \\ \hline \end{array} \cdot \begin{array}{|c|c|} \hline 2 & 5 \\ \hline 0 & 1 \\ \hline 3 & 4 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 3 & 6 \\ \hline 18 & 32 \\ \hline \end{array}$$

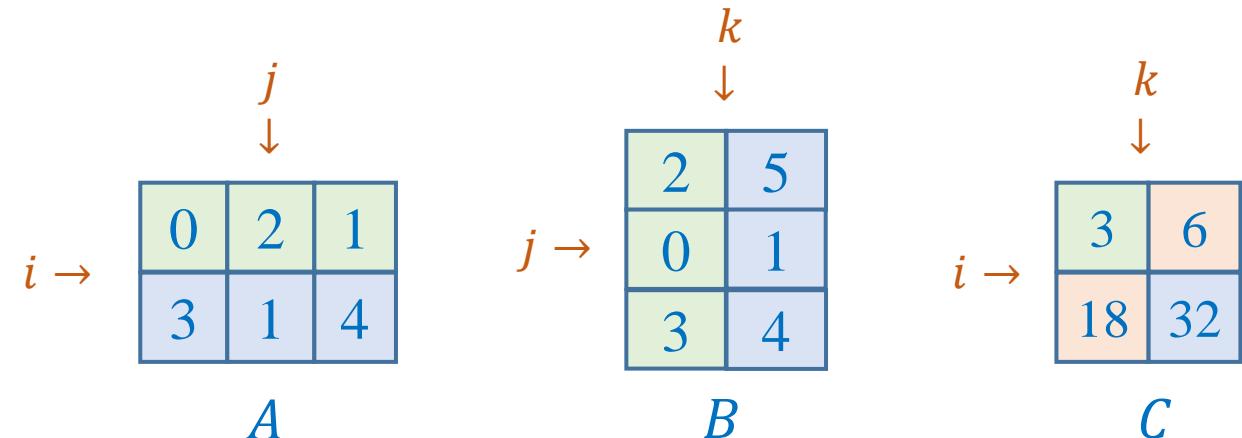
A *B* *C*

```
import torch

A = torch.Tensor([[0, 2, 1],
                  [3, 1, 4]])
B = torch.Tensor([[2, 5],
                  [0, 1],
                  [3, 4]])
C = torch.einsum('ij,jk->ik', A, B)
print(C)

tensor([[ 3.,  6.],
        [18., 32.]])
```

ij, jk → *ik*



$$C_{ik} = \sum_j A_{ij} B_{jk} = A_{i1} B_{1k} + A_{i2} B_{2k} + \dots$$

Einstein Summation

Matrix multiplication

$$\begin{array}{|c|c|c|} \hline 0 & 2 & 1 \\ \hline 3 & 1 & 4 \\ \hline \end{array} \cdot \begin{array}{|c|c|} \hline 2 & 5 \\ \hline 0 & 1 \\ \hline 3 & 4 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 3 & 6 \\ \hline 18 & 32 \\ \hline \end{array}$$

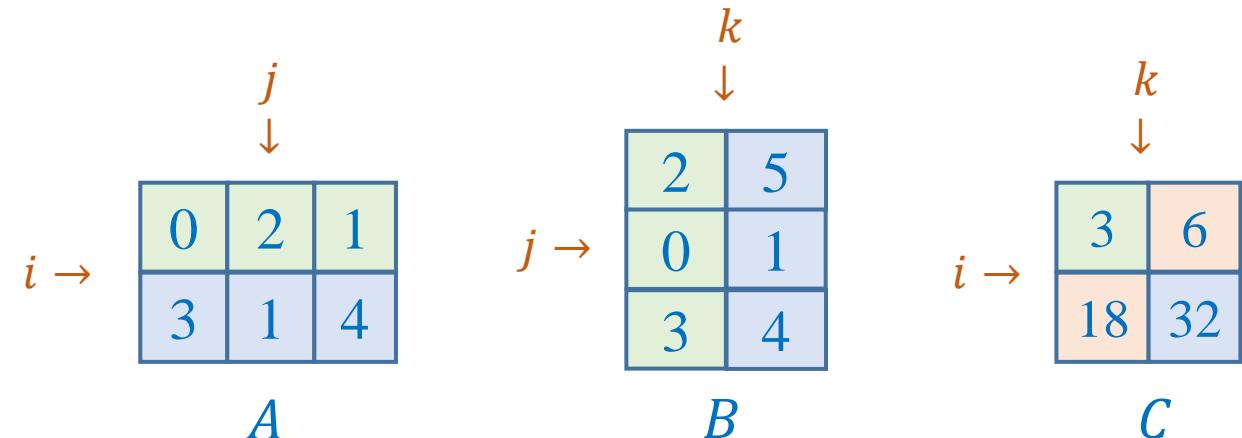
A *B* *C*

```
import torch

A = torch.Tensor([[0, 2, 1],
                  [3, 1, 4]])
B = torch.Tensor([[2, 5],
                  [0, 1],
                  [3, 4]])
C = torch.einsum('ij,jk', A, B)
print(C)

tensor([[ 3.,  6.],
        [18., 32.]])
```

ij, jk →



$$C_{ik} = \sum_j A_{ij} B_{jk} = A_{i1} B_{1k} + A_{i2} B_{2k} + \dots$$

Repeated indices are summed if the output indices are not specified.

QUIZ TIME

Outline

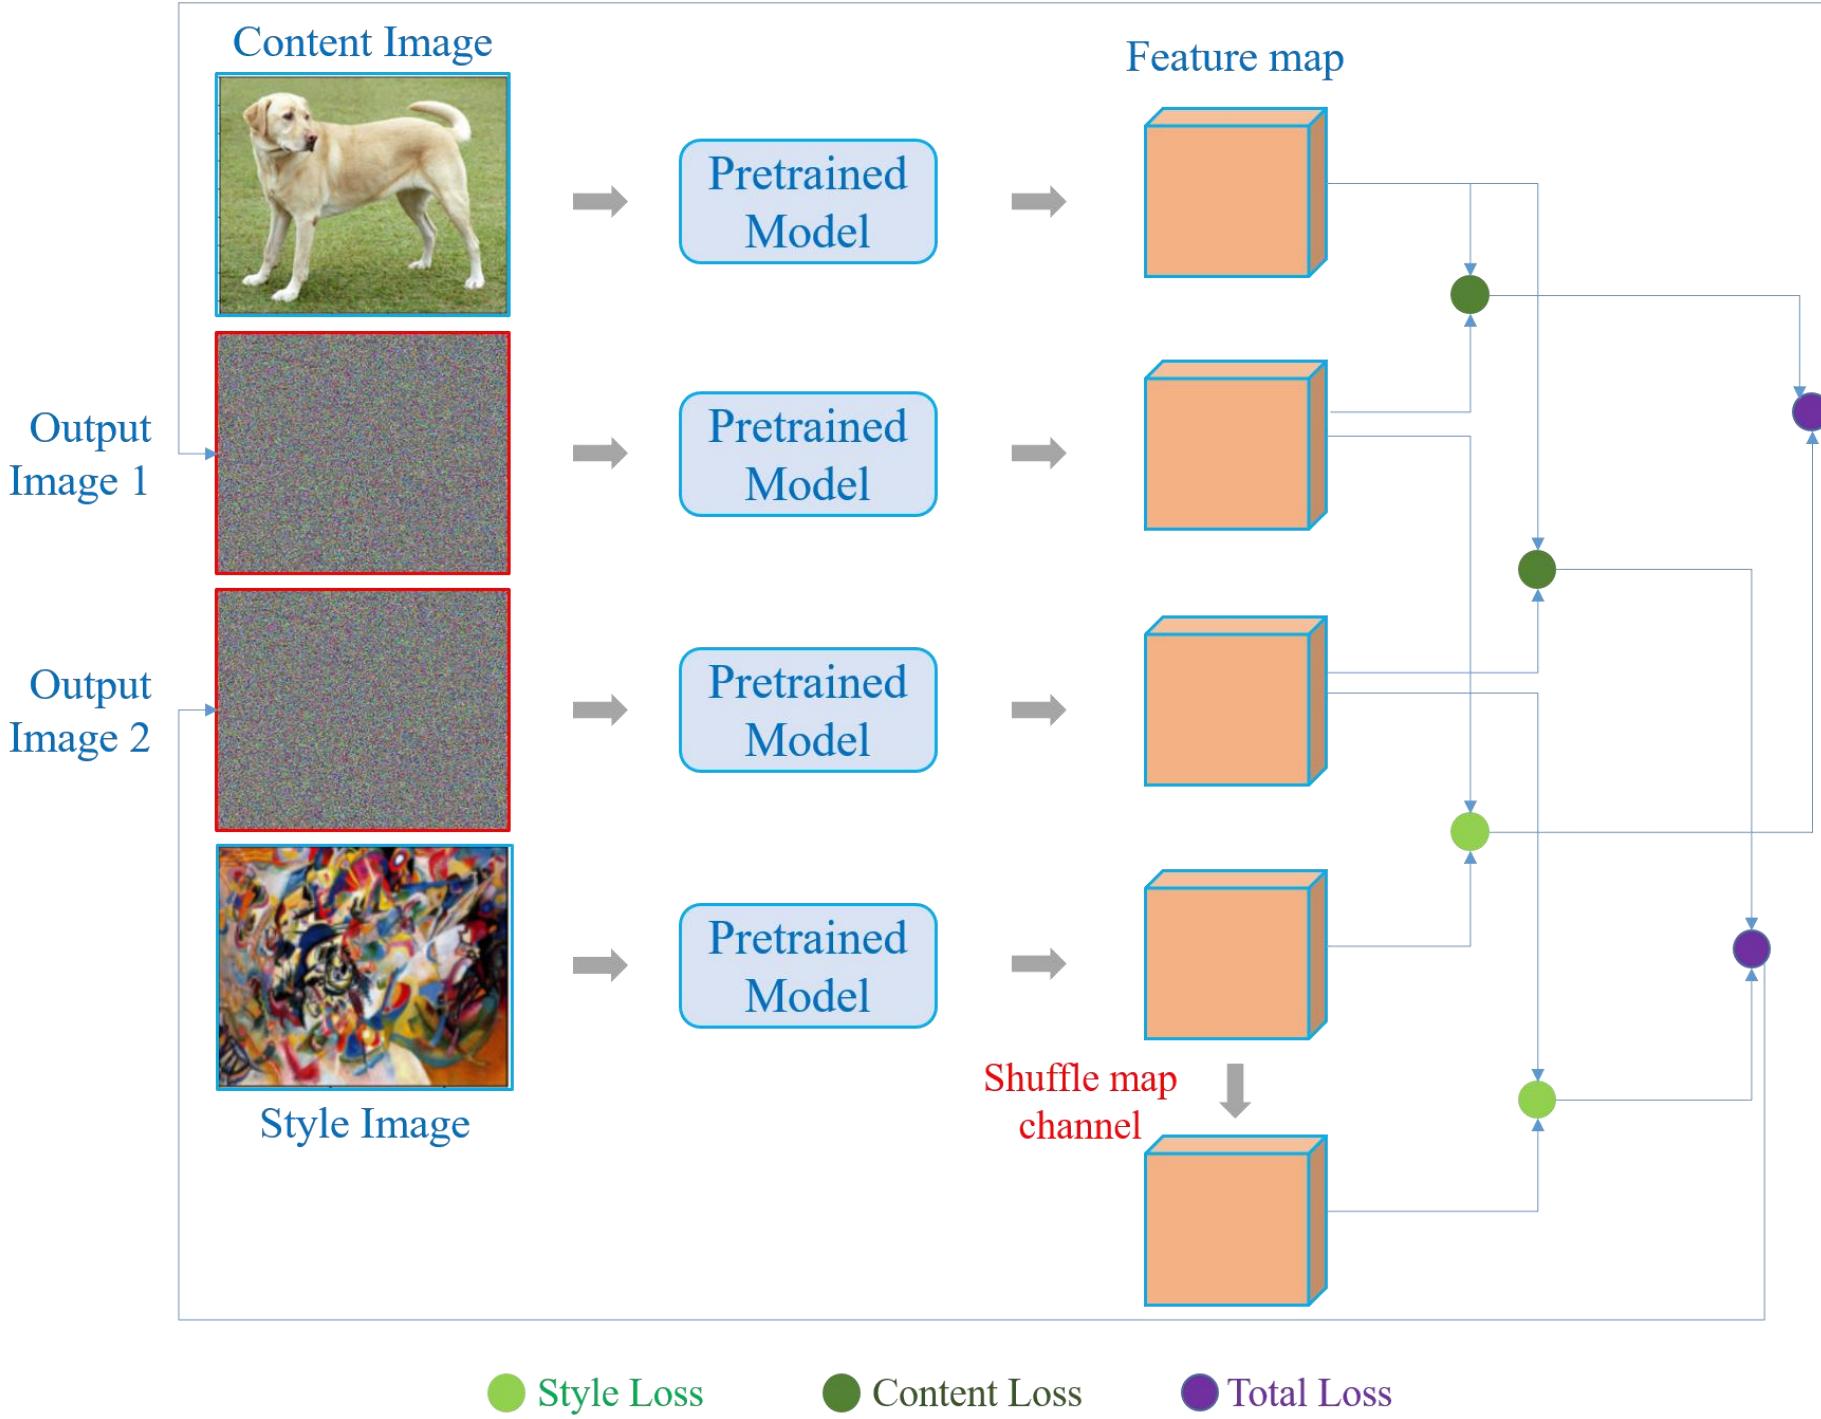
- Quick Review
- Content Loss+Style Loss
- Implementation
- Einstein Summation
- Extensions

Multi-modal Style Transfer



Multi-modal
Style Transfer



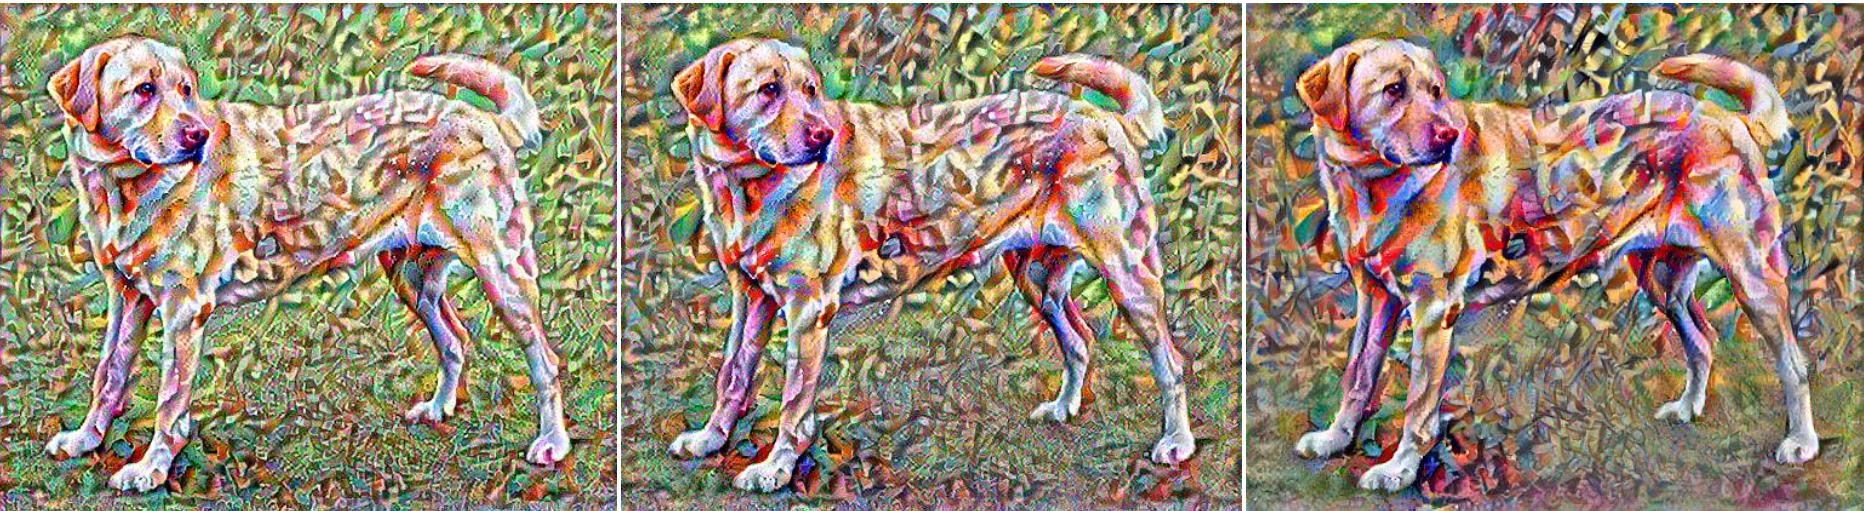


Multi-modal Style Transfer

❖ Shuffle feature maps

```
1 style_outputs_1 = self.vgg(preprocessed_input)
2 style_outputs_2 = [tf.roll(style_output,
3                                         shift=[style_output.shape[1]//2,
4                                         style_output.shape[2]//2,
5                                         style_output.shape[3]//2],
6                                         axis=[1,2,3])
7                                         for style_output in style_outputs_1]
```

Output
Image 1

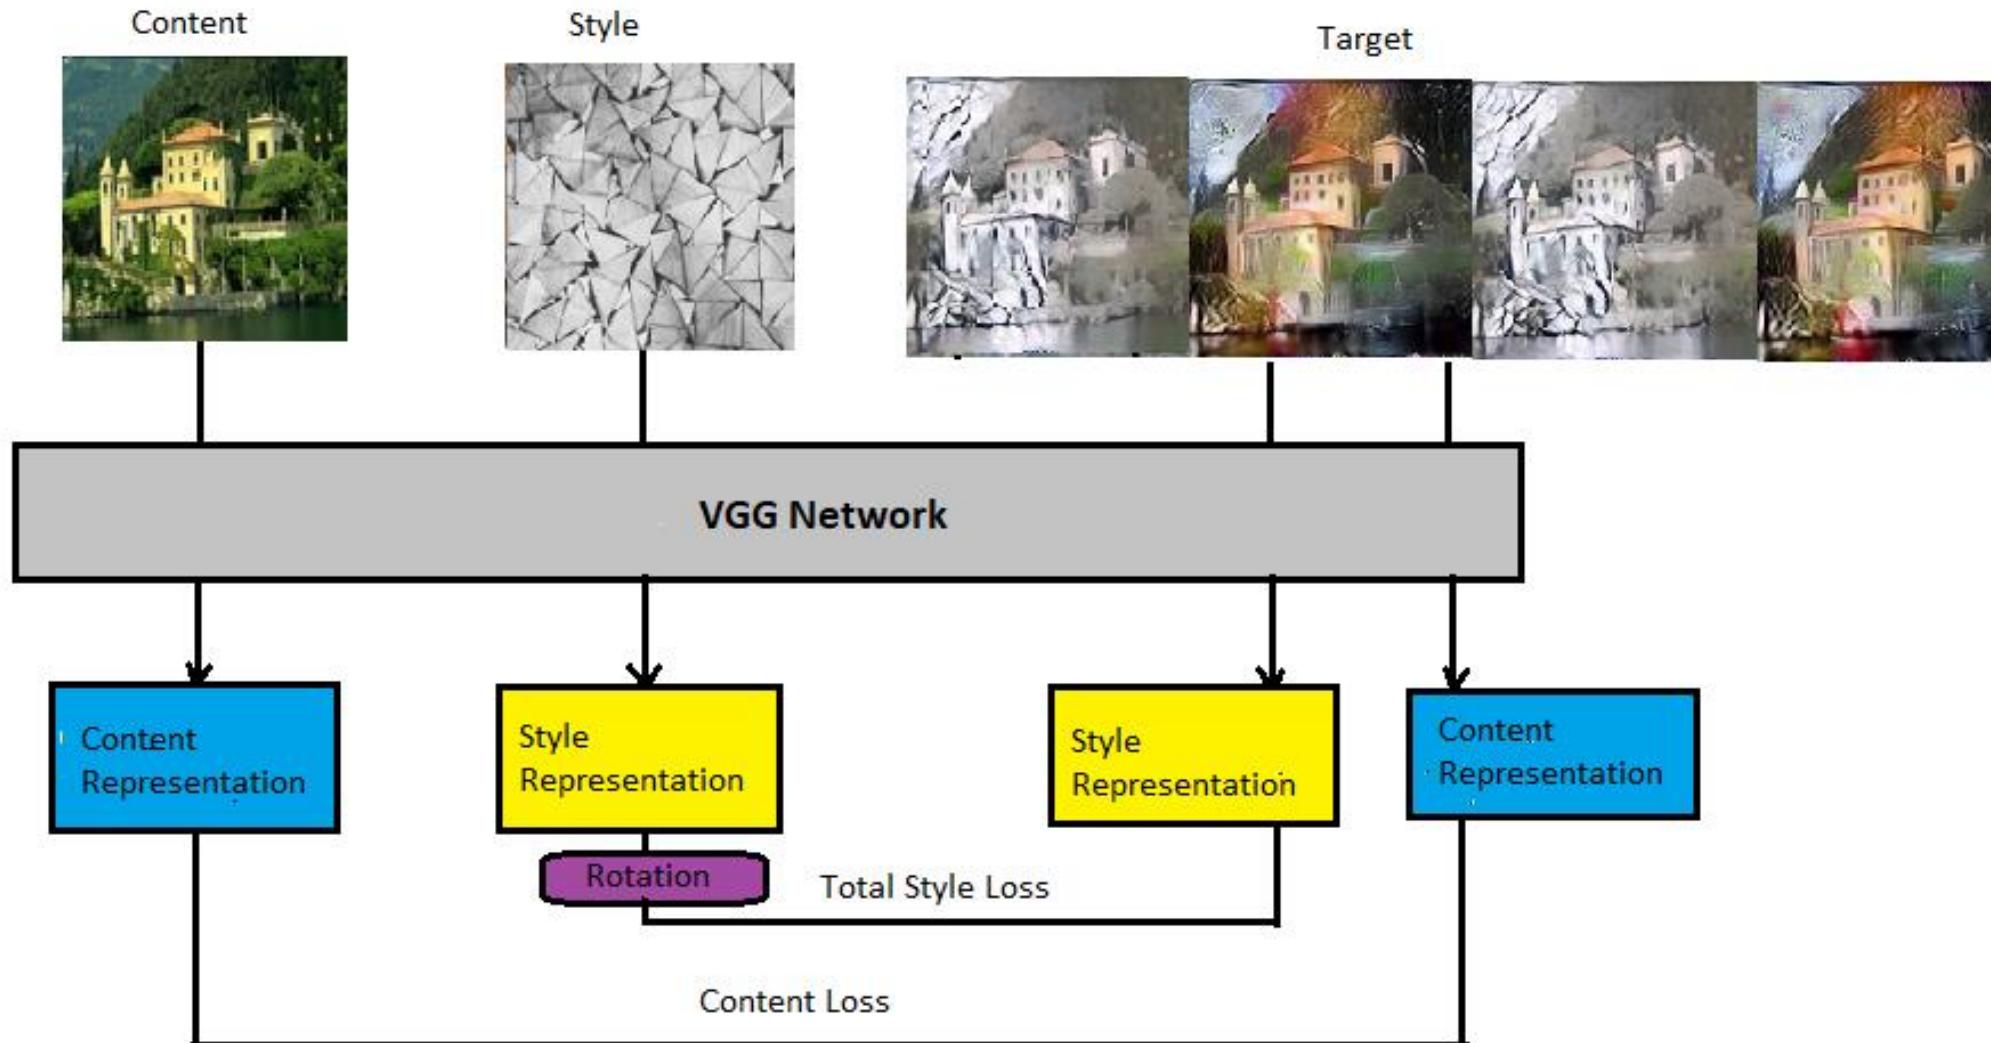


Output
Image 2

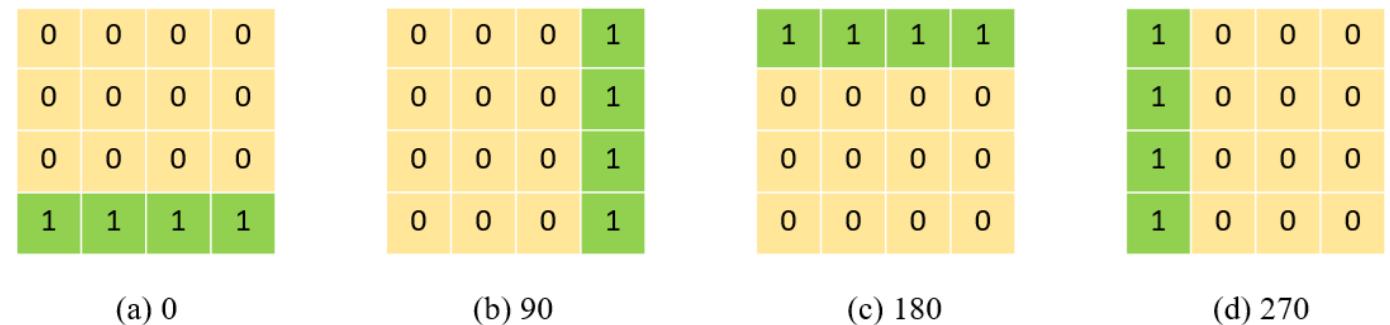


Multi-modal Style Transfer

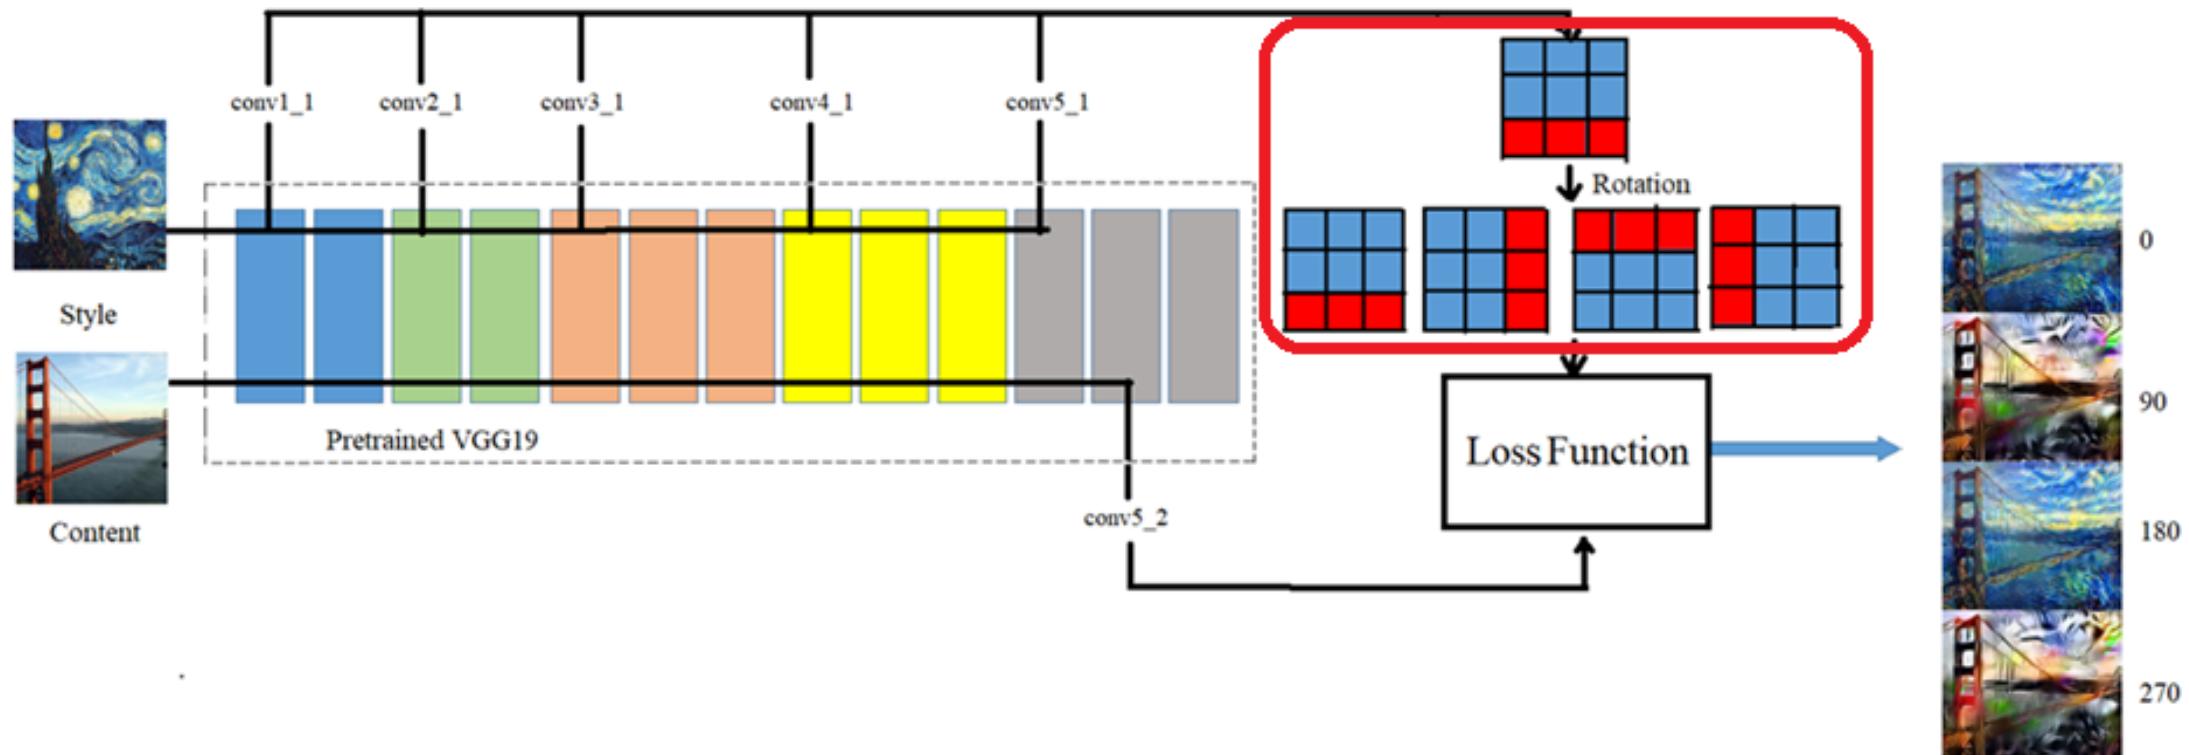
❖ Deep feature rotation



Multi-modal Style Transfer

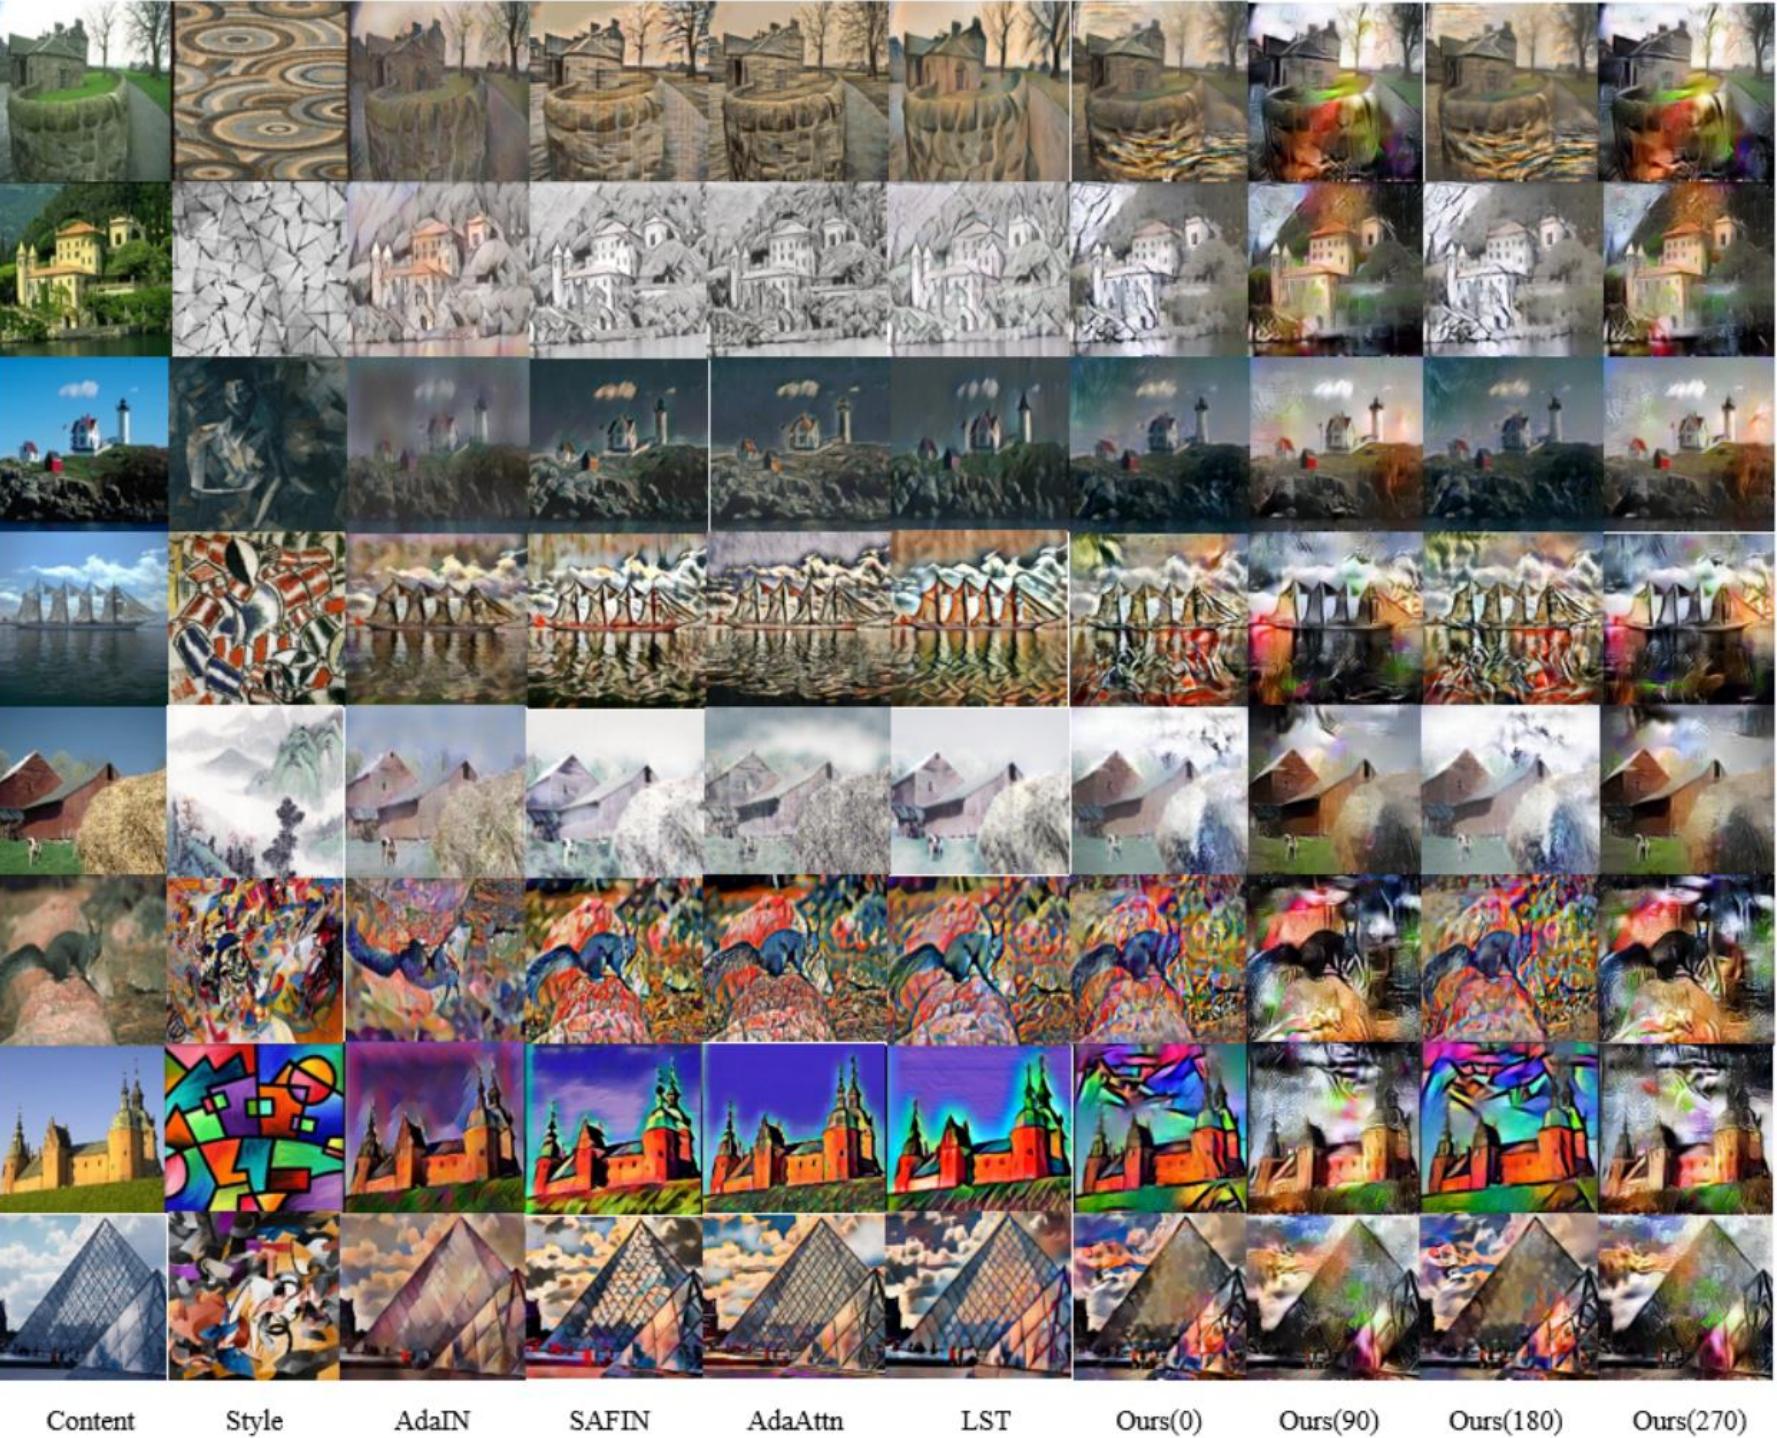


❖ Deep feature rotation



Multi-modal Style Transfer

❖ Results



❖ Two style images

7.2.Style_transfer_1content_2style.ipynb



Multi-modal Style Transfer

❖ Two style images

7.3.Style_transfer_1content_2style.ipynb



Epoch 1



Epoch 3

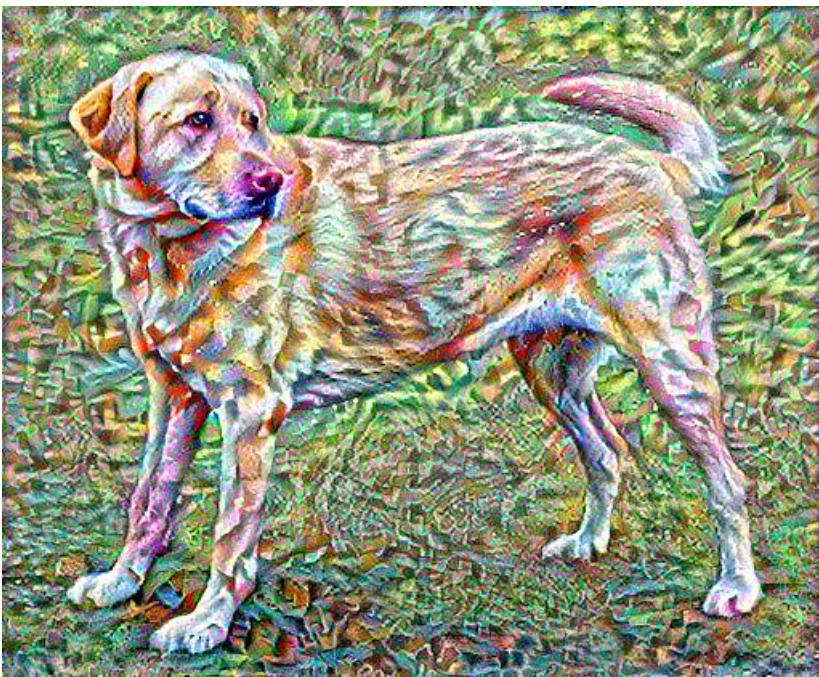


Epoch 10

Multi-modal Style Transfer

❖ Two style images

7.4.Style_transfer_1content_2style.ipynb



Epoch 1



Epoch 3



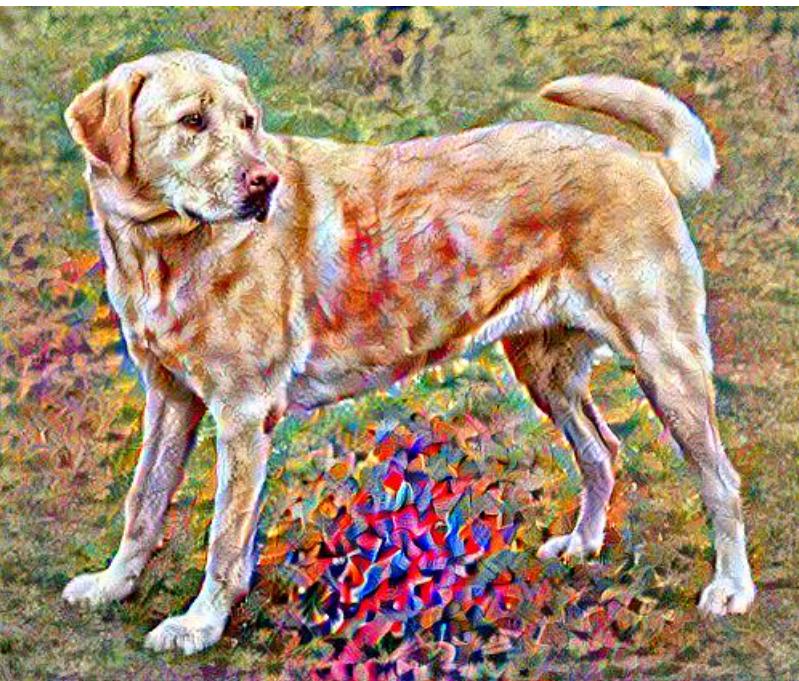
Epoch 10

Multi-modal Style Transfer

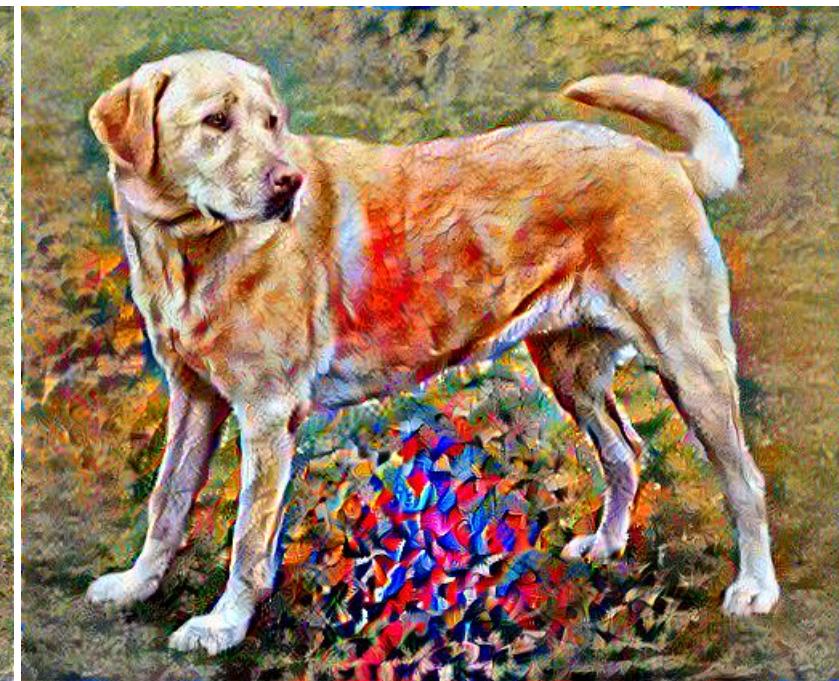
❖ Using the Euclidian Distance



Epoch 1



Epoch 3



Epoch 10

Multi-modal Style Transfer

❖ Different measures

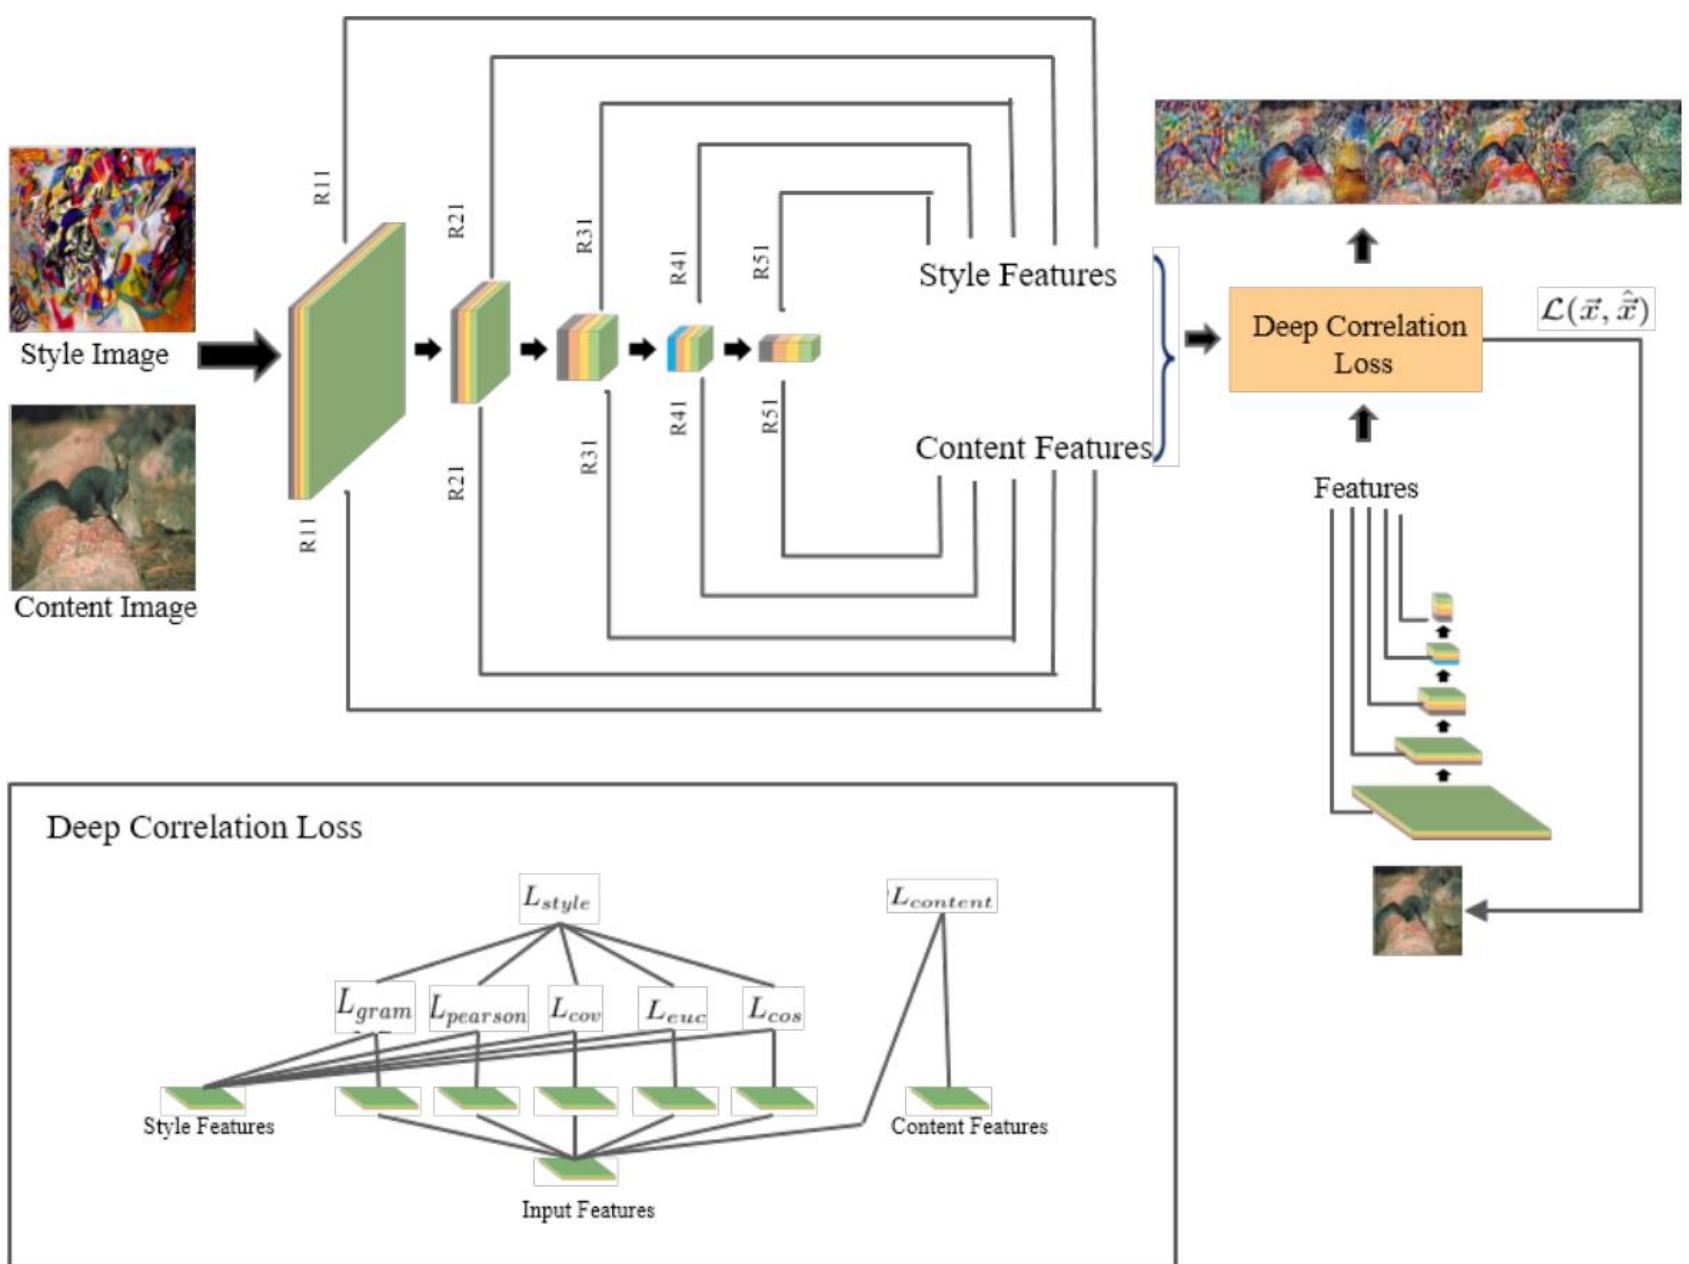


FIGURE 1: Overview of DeCorMST and details of its sub-networks. The architecture consists of a backbone of VGG extracting features from style, content, and generated images. These feature maps go through a deep correlation loss layer that integrates losses based on different correlations.

Multi-modal Style Transfer

❖ Results

Deep Correlation
Multimodal Style Transfer



(a) Content Image

(b) Style Image

(c) Gatys's method

(d) SAFIN



(e) Output 1

(f) Output 2

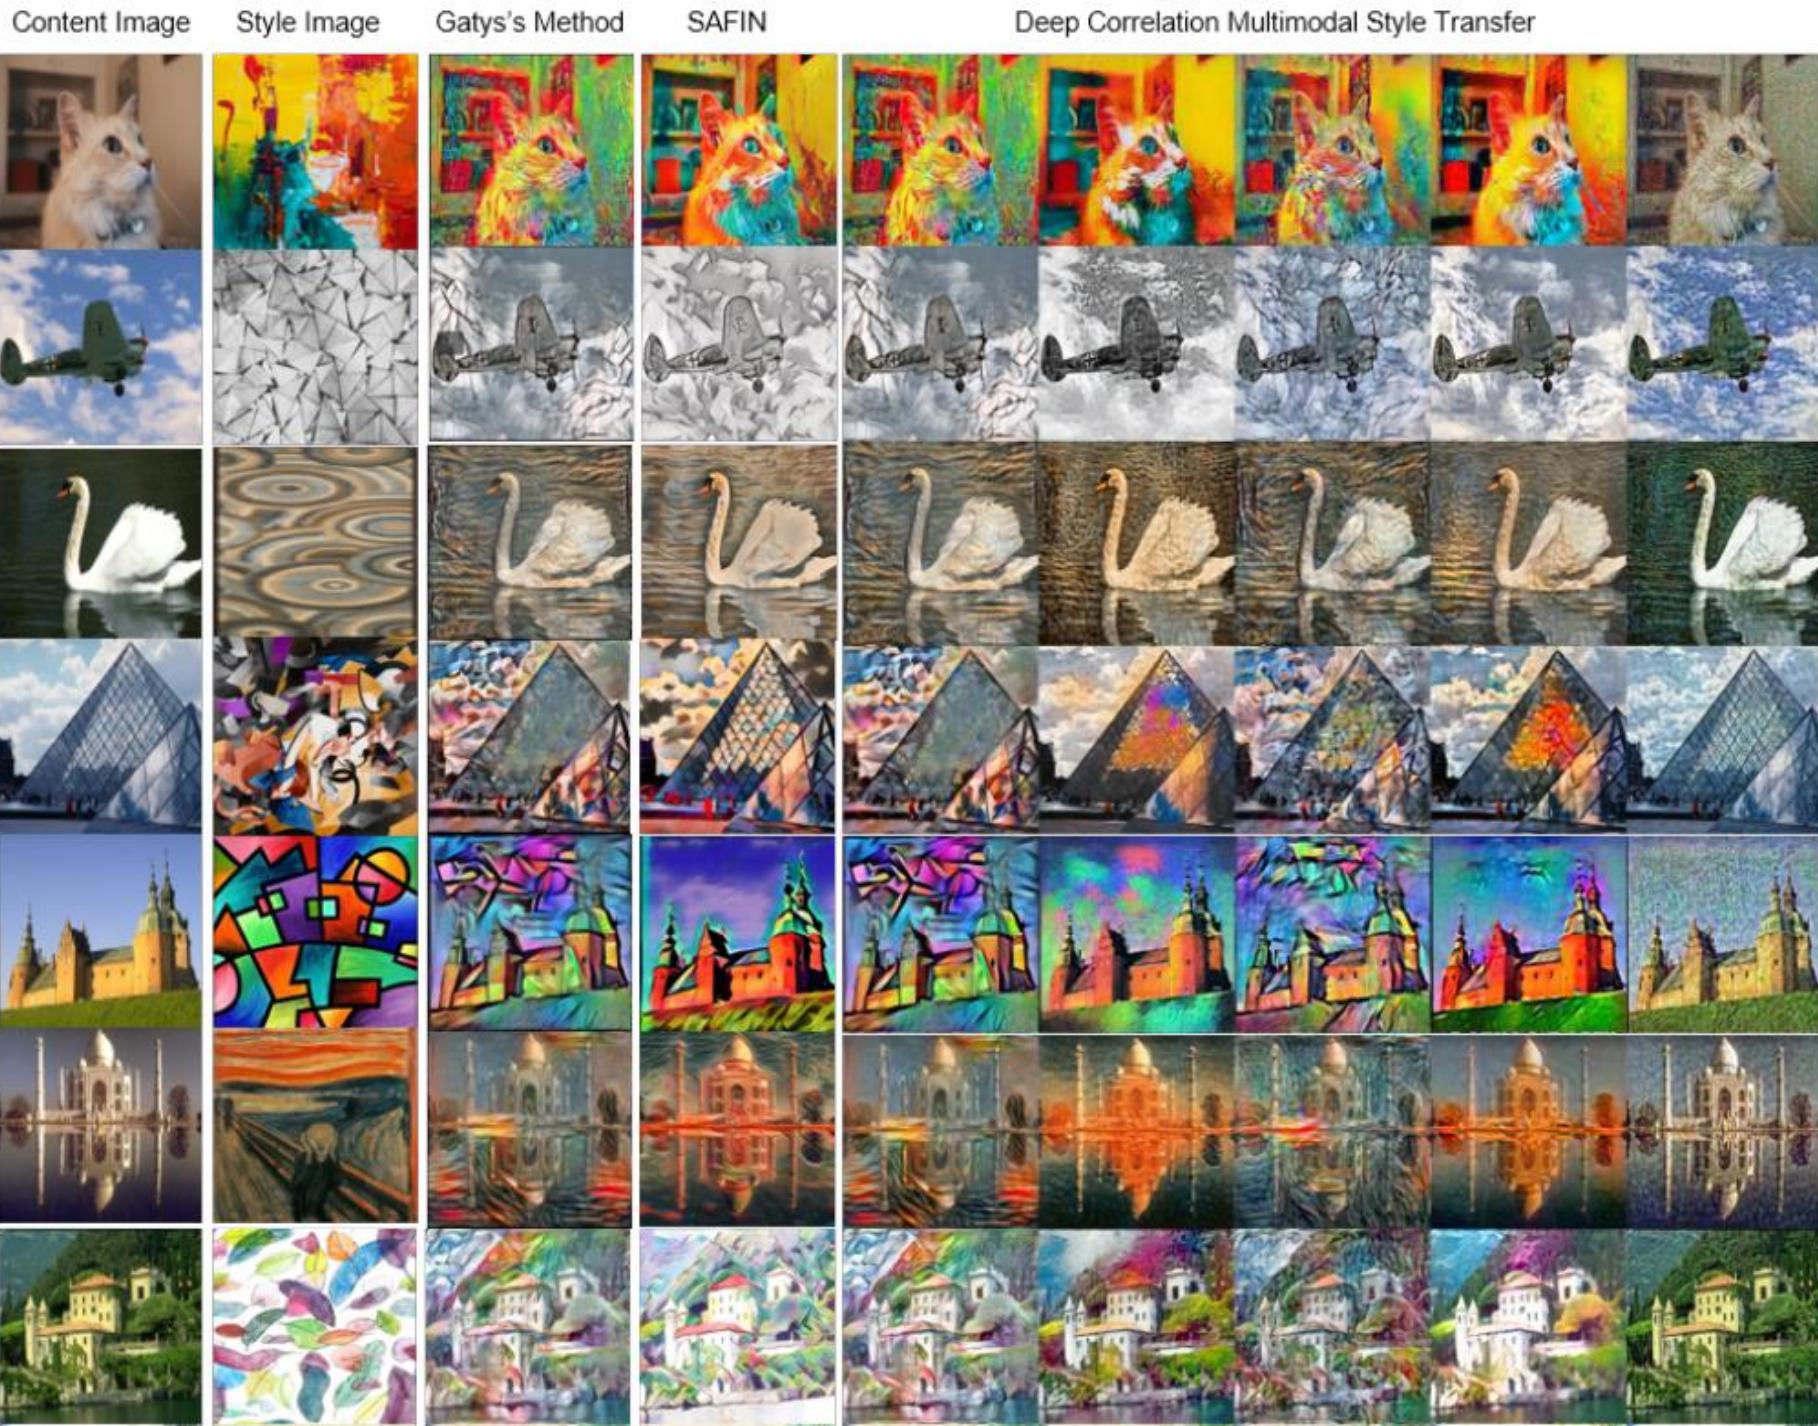
(g) Output 3

(h) Output 4

(i) Output 5

Multi-modal Style Transfer

❖ Results

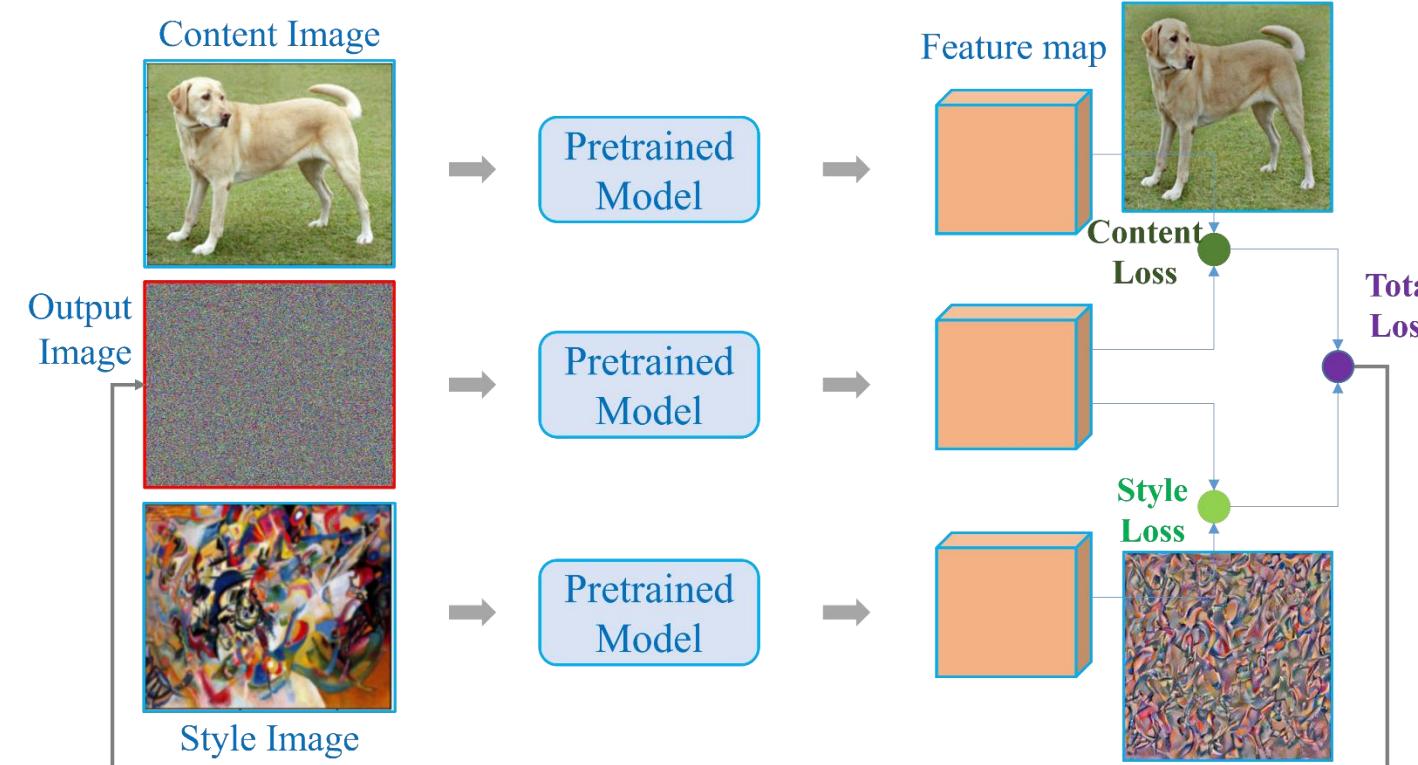


Summary

- Quick Review
- Content Loss+Style Loss
- Implementation
- Einstein Summation
- Extensions

Summary

- ✓ Studied about content loss + style loss
- ✓ Implemented simple methods
- ✓ Studied the Einsum and applied to style transfer
- ✓ Discussed ideas beyond the traditional style transfer



```
import torch
```

```
x = torch.Tensor([[4, 0, 3, 1],  
[3, 0, 1, 2],  
[4, 5, 0, 3]])  
G = torch.einsum('ij,jk->ik', x, x.t())  
print(G)
```

```
tensor([[26., 17., 19.],  
[17., 14., 18.],  
[19., 18., 50.]])
```



